

# Deep Reinforcement Learning

Overview of main articles

Part 1. Value-based algorithms

---

Sergey Ivanov

November 1, 2018

MSU

# Table of contents i

Reinforcement Learning [reminder]

Deep Q-learning (2014)

Stabilizing Q-learning

Target-network heuristic (2015)

Double DQN (2015)

Dueling DQN (2016)

Prioritized replay memory (2015)

Noisy networks for exploration (2017)

Categorical DQN (2017)

Rainbow DQN (2018)

# Reinforcement Learning

## [reminder]

---

MDP is  $\{\mathbb{S}, \mathbb{A}, \mathbb{T}, r\}$ :

$\mathbb{S}$  — set of states

$\mathbb{A}$  — set of actions

$\mathbb{T}$  — probability  $p(s' \mid s, a)$ , where  $s, s' \in \mathbb{S}, a \in \mathbb{A}$

$r$  — function  $\mathbb{S} \rightarrow \mathbb{R}$

We search for policy  $\pi : \mathbb{S} \rightarrow \mathbb{A}$  which maximizes<sup>1</sup>

$$\mathbb{E} \sum_t r(s_t)$$

---

<sup>1</sup>over what probability distributions is this expectation?

We search for policy  $\pi : \mathbb{S} \rightarrow \mathbb{A}$  which maximizes<sup>1</sup>

$$\mathbb{E} \sum_t r(s_t)$$

This goal does not suit infinite horizon case, so for generalization purposes goal is substituted with

$$\mathbb{E} \sum_t \gamma^t r(s_t)$$

for  $\gamma \in (0, 1)$ .

---

<sup>1</sup>over what probability distributions is this expectation?

For convenience<sup>2</sup>:

$$R = \sum_t \gamma^t r(s_t)$$

---

<sup>2</sup>*What does it depend on?*

# Definitions

For convenience<sup>2</sup>:

$$R = \sum_t \gamma^t r(s_t)$$

For **given policy**  $\pi$ :

$$V^\pi(s) = \mathbb{E}R \mid s_0 = s$$

$$Q^\pi(s, a) = \mathbb{E}V(s') \mid s, a$$

---

<sup>2</sup>What does it depend on?



# Definitions

For convenience<sup>2</sup>:

$$R = \sum_t \gamma^t r(s_t)$$

For **given policy**  $\pi$ :

$$V^\pi(s) = \mathbb{E}R \mid s_0 = s$$

$$Q^\pi(s, a) = \mathbb{E}V(s') \mid s, a$$

Let  $\pi^*$  be optimal policy.

---

<sup>2</sup>What does it depend on?

# Bellman Equation

For every  $\pi$  it's true:

$$Q^\pi(s, a) = \mathbb{E} [r(s') + Q^\pi(s', \pi(s'))]$$

# Bellman Equation

For every  $\pi$  it's true:

$$Q^{\pi}(s, a) = \mathbb{E} [r(s') + Q^{\pi}(s', \pi(s'))]$$

It's also true for  $\pi^*$ :

$$Q^{\pi^*}(s, a) = \mathbb{E} [r(s') + Q^{\pi^*}(s', \pi^*(s'))] \quad (1)$$

# Bellman Equation

For every  $\pi$  it's true:

$$Q^{\pi}(s, a) = \mathbb{E} [r(s') + Q^{\pi}(s', \pi(s'))]$$

It's also true for  $\pi^*$ :

$$Q^{\pi^*}(s, a) = \mathbb{E} [r(s') + Q^{\pi^*}(s', \pi^*(s'))] \quad (1)$$

Note:

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q^{\pi^*}(s, a) \quad (2)$$

# Bellman Equation

For every  $\pi$  it's true:

$$Q^{\pi}(s, a) = \mathbb{E} [r(s') + Q^{\pi}(s', \pi(s'))]$$

It's also true for  $\pi^*$ :

$$Q^{\pi^*}(s, a) = \mathbb{E} [r(s') + Q^{\pi^*}(s', \pi^*(s'))] \quad (1)$$

Note:

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q^{\pi^*}(s, a) \quad (2)$$

# Bellman Equation

For every  $\pi$  it's true:

$$Q^{\pi}(s, a) = \mathbb{E} [r(s') + Q^{\pi}(s', \pi(s'))]$$

It's also true for  $\pi^*$ :

$$Q^{\pi^*}(s, a) = \mathbb{E} [r(s') + Q^{\pi^*}(s', \pi^*(s'))] \quad (1)$$

Note:

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q^{\pi^*}(s, a) \quad (2)$$

Insert (2) into (1):

## Bellman Equation

$$Q^{\pi^*}(s, a) = \mathbb{E} \left[ r(s') + Q^{\pi^*}(s', \underset{a}{\operatorname{argmax}} Q^{\pi^*}(s', a)) \right]$$

# Bellman Equation

For every  $\pi$  it's true:

$$Q^\pi(s, a) = \mathbb{E} [r(s') + Q^\pi(s', \pi(s'))]$$

It's also true for  $\pi^*$ :

$$Q^{\pi^*}(s, a) = \mathbb{E} [r(s') + Q^{\pi^*}(s', \pi^*(s'))] \quad (1)$$

Note:

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q^{\pi^*}(s, a) \quad (2)$$

Insert (2) into (1):

## Bellman Equation

$$Q^{\pi^*}(s, a) = \mathbb{E} [r(s') + \max_a Q^{\pi^*}(s', a)]$$

# Temporal Difference Learning

For **finite-state case**  $Q^{\pi^*}$  is finite vector of unknown values.

Bellman equations can be solved using point iteration:

$$Q_{t+1}(s, a) = \mathbb{E} \left[ r(s') + \max_a Q_t(s', a) \right]$$



# Temporal Difference Learning

For **finite-state case**  $Q^{\pi^*}$  is finite vector of unknown values.

Bellman equations can be solved using point iteration:

$$Q_{t+1}(s, a) = \mathbb{E} \left[ r(s') + \max_a Q_t(s', a) \right]$$

**Problem:** expectation.

# Temporal Difference Learning

For **finite-state case**  $Q^{\pi^*}$  is finite vector of unknown values.

Bellman equations can be solved using point iteration:

$$Q_{t+1}(s, a) = \mathbb{E} \left[ r(s') + \max_a Q_t(s', a) \right]$$

**Problem:** expectation.

## Temporal Difference Learning

$$Q_{t+1}(s, a) = \alpha Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) \right]$$

# Temporal Difference Learning

For **finite-state case**  $Q^{\pi^*}$  is finite vector of unknown values.

Bellman equations can be solved using point iteration:

$$Q_{t+1}(s, a) = \mathbb{E} \left[ r(s') + \max_a Q_t(s', a) \right]$$

**Problem:** expectation.

## Temporal Difference Learning


$$Q_{t+1}(s, a) = \alpha Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) \right]$$

✓ Is a *contraction mapping*  $\Rightarrow$  converges.

# Deep Q-learning (2014)


---

- \* No prepared features for each game.
- \* Screen image as input.
- \* Finite-state case... not quite finite.



Atari games

- \* No prepared features for each game.
- \* Screen image as input.
- \* Finite-state case... not quite finite.




Atari games





We want to approximate  $Q(s, a)$  with neural net.

# Q-network



# Q-network

Option 1  
Requires forward pass for each action<sup>1</sup>



Option 2  
Number of actions must be adequate



---

<sup>1</sup>Is there a case when option 1 might be better?



# Q-network

Option 1  
Requires forward pass for each action<sup>1</sup>



Option 2  
Number of actions must be adequate

Atari: up to 18 discrete actions. Use option 2.

---

<sup>1</sup>Is there a case when option 1 might be better?

# TD-learning to gradient descent

TD-learning is «similar» to gradient descent.

$$Q_{t+1}(s, a) = \alpha Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) \right]$$

# TD-learning to gradient descent

TD-learning is «similar» to gradient descent.

$$\begin{aligned} Q_{t+1}(s, a) &= \alpha Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) \right] = \\ &= Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) - Q_t(s, a) \right] \end{aligned}$$

# TD-learning to gradient descent

TD-learning is «similar» to gradient descent.

$$\begin{aligned} Q_{t+1}(s, a) &= \alpha Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) \right] = \\ &= Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) - Q_t(s, a) \right] = \\ &= Q_t(s, a) - \eta \nabla_Q L \end{aligned}$$

# TD-learning to gradient descent

TD-learning is «similar» to gradient descent.

$$\begin{aligned} Q_{t+1}(s, a) &= \alpha Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) \right] = \\ &= Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) - Q_t(s, a) \right] = \\ &= Q_t(s, a) - \eta \nabla_Q L \end{aligned}$$

# TD-learning to gradient descent

TD-learning is «similar» to gradient descent.

$$\begin{aligned} Q_{t+1}(s, a) &= \alpha Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) \right] = \\ &= Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) - Q_t(s, a) \right] = \\ &= Q_t(s, a) - \eta \nabla_Q L \end{aligned}$$

Let  $y = r(s') + \max_a Q_t(s', a)$ .

# TD-learning to gradient descent

TD-learning is «similar» to gradient descent.

$$\begin{aligned} Q_{t+1}(s, a) &= \alpha Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) \right] = \\ &= Q_t(s, a) + (1 - \alpha) \left[ r(s') + \max_a Q_t(s', a) - Q_t(s, a) \right] = \\ &= Q_t(s, a) - \eta \nabla_Q L \end{aligned}$$

Let  $y = r(s') + \max_a Q_t(s', a)$ .

If dependence of  $y$  from  $Q$  is ignored:

$$L = (Q_t(s, a) - y)^2$$

With  $Q(s, a)$  as neural net, its parameters  $\theta$  determine function.

$$Q_{t+1}(s, a, \theta) = Q_t(s, a, \theta) - \eta \nabla_Q L$$



With  $Q(s, a)$  as neural net, its parameters  $\theta$  determine function.

$$Q_{t+1}(s, a, \theta) = Q_t(s, a, \theta) - \eta \nabla_Q L$$

Let's move gradient descent from space of  $Q$  functions to  $\theta$ !

$$\theta_{t+1} = \theta_t - \beta \nabla_{\theta} L$$

With  $Q(s, a)$  as neural net, its parameters  $\theta$  determine function.

$$Q_{t+1}(s, a, \theta) = Q_t(s, a, \theta) - \eta \nabla_Q L$$

Let's move gradient descent from space of  $Q$  functions to  $\theta$ !

$$\theta_{t+1} = \theta_t - \beta \nabla_{\theta} L$$

**Problems:**

× `batch_size = 1`. Wow.

With  $Q(s, a)$  as neural net, its parameters  $\theta$  determine function.

$$Q_{t+1}(s, a, \theta) = Q_t(s, a, \theta) - \eta \nabla_Q L$$

Let's move gradient descent from space of  $Q$  functions to  $\theta$ !

$$\theta_{t+1} = \theta_t - \beta \nabla_{\theta} L$$

## Problems:

- × `batch_size = 1`. Wow.
- × Target  $y$  changes after each step.

With  $Q(s, a)$  as neural net, its parameters  $\theta$  determine function.

$$Q_{t+1}(s, a, \theta) = Q_t(s, a, \theta) - \eta \nabla_Q L$$

Let's move gradient descent from space of  $Q$  functions to  $\theta$ !

$$\theta_{t+1} = \theta_t - \beta \nabla_{\theta} L$$

## Problems:

- × `batch_size = 1`. Wow.
- × Target  $y$  changes after each step.
- × All theoretical guarantees are lost.



Utilize all experienced transitions  $(s, a, s', r, done)$  for generating a batch for stochastic optimization step.



Utilize all experienced transitions  $(s, a, s', r, done)$  for generating a batch for stochastic optimization step.

Pretend on each step that loss function is

$$\mathbb{E}_{(s,a,s',r,done)}(Q(s,a,\theta) - y(s',r,done))^2$$

Batch of transitions is sampled uniformly from memory.



Utilize all experienced transitions  $(s, a, s', r, done)$  for generating a batch for stochastic optimization step.

Pretend on each step that loss function is

$$\mathbb{E}_{(s,a,s',r,done)}(Q(s,a,\theta) - y(s',r,done))^2$$

Batch of transitions is sampled uniformly from memory.

- ✓ Decorrelates samples.
- \* Target  $y$  can be calculated only for this batch.
- \* Only last  $N$  observed transitions may be stored

**Problem:** at the very beginning trajectories generated by  $\pi(s) = \underset{a}{\operatorname{argmax}} Q(s, a, \theta)$  are very similar.



**Problem:** at the very beginning trajectories generated by  $\pi(s) = \underset{a}{\operatorname{argmax}} Q(s, a, \theta)$  are very similar.



Choose random actions sometimes.

For example, with probability  $\epsilon$ .

$\epsilon$  should be big at the beginning and small at the end.

$\varepsilon$  should be big at the beginning and small at the end.

Atari:  $\varepsilon(i) = 0.01 + 0.99 \exp\{-\frac{i}{30000}\}$  where  $i$  is frames counter.

- Gray-scale frames were downsampled and cropped to 84x84.
- Last 4 frames<sup>3</sup> were considered as state to satisfy MDP Markov's property.
- Same NN architecture was used for all games: 3 convolutional<sup>4</sup> and 2 feedforward layers.

---

<sup>3</sup>3 for Space Invaders cause of laser blinking period

<sup>4</sup>why no max pooling here?

### Playing Atari with Deep Reinforcement Learning (2014)

- Reward was restricted to  $\{+1, 0, -1\}$ . Allowed to use same learning rate for all games.
- :( 50 hours per game / 10 000 000 frames per game.
- :} Bought by Google after 7 games.

# Stabilizing Q-learning

---

Recall our target on each step:

$$y(s', r) = r + \max_{a'} Q(s', a', \theta)$$

- Changes each frame
- Formally depends on  $\theta$
- "Correlates" with actions chosen during playing
- Tends to overestimate true  $V(s')$

⇒ loss is completely unstable and can even diverge.

## Target network (2015)



Change the target not every step, but each  $K$ -th step.



## Target network (2015)



Change the target not every step, but each  $K$ -th step.

For this purpose:



- Make a copy of Q-network, *target network*, with parameters  $\theta^-$
- Use it on every step to calculate

$$y(s', r) = r + \max_{a'} Q^{\text{target}}(s', a', \theta^-)$$

- Each  $K$ -th step update  $\theta^-$  with current Q-network's weights  $\theta$ .

# Can be seen on loss

✓ Loss really stabilized!



# Value overestimation

Recall our target is proxy of  $V^{\pi^*}(s', a')$

$$y(s', r) = r + \max_{a'} Q(s', a', \theta)$$

**Practice:** this proxy overestimates true value of states.

**Intuition:** this max operator will prefer actions, for which  $Q(s', a', \theta)$  is overestimating true value due to approximation or luck.

# Action Selection vs Evaluation

Recall Bellman Equation derivation and untangle our target:

$$y(s', r) = r + \max_{a'} Q(s', a', \theta) = r + Q(s', \underset{a'}{\operatorname{argmax}} Q(s', a', \theta), \theta)$$

# Action Selection vs Evaluation

Recall [Bellman Equation derivation](#) and untangle our target:

$$y(s', r) = r + \max_{a'} Q(s', a', \theta) = r + Q(s', \underset{a'}{\operatorname{argmax}} Q(s', a', \theta), \theta)$$

- \*  $\underset{a'}{\operatorname{argmax}} Q(s', a', \theta)$  is *action selection*
- \*  $Q(s', a', \theta)$  is *action evaluation*

# Action Selection vs Evaluation

Recall [Bellman Equation derivation](#) and untangle our target:

$$y(s', r) = r + \max_{a'} Q(s', a', \theta) = r + Q(s', \underset{a'}{\operatorname{argmax}} Q(s', a', \theta), \theta)$$

- \*  $\underset{a'}{\operatorname{argmax}} Q(s', a', \theta)$  is *action selection*
- \*  $Q(s', a', \theta)$  is *action evaluation*

General idea:



Use different approximations for evaluation and for selection to avoid *max*.

# Two Q-learnings

## Basic way to do this:

run two Q-learning algorithms with two approximations of  $Q^{\pi^*}$ :  
 $Q_1(s, a, \theta_1)$  and  $Q_2(s, a, \theta_2)$ .

# Two Q-learnings

## Basic way to do this:

run two Q-learning algorithms with two approximations of  $Q^{\pi^*}$ :  
 $Q_1(s, a, \theta_1)$  and  $Q_2(s, a, \theta_2)$ .

Targets for Q-learnings:

$$y_1 = r + Q_2(s', \underset{a'}{\operatorname{argmax}} Q_1(s', a', \theta_1), \theta_2)$$

$$y_2 = r + Q_1(s', \underset{a'}{\operatorname{argmax}} Q_2(s', a', \theta_2), \theta_1)$$



# Double DQN (2015)

## Deep Reinforcement Learning with Double Q-learning (2015)

- **more convenient way to do this:**



Use target network as one of two approximations.

---

<sup>5</sup>how many backwards?

# Double DQN (2015)

## Deep Reinforcement Learning with Double Q-learning (2015)

- more convenient way to do this:



Use target network as one of two approximations.

$$y = r + Q^{\text{target}}(s', \underset{a'}{\operatorname{argmax}} Q(s', a', \theta), \theta^-)$$

---

<sup>5</sup>how many backwards?

# Double DQN (2015)

## Deep Reinforcement Learning with Double Q-learning (2015)

- more convenient way to do this:



Use target network as one of two approximations.

$$y = r + Q^{\text{target}}(s', \underset{a'}{\operatorname{argmax}} Q(s', a', \theta), \theta^-)$$

- \* Keep ignoring dependence of  $y$  from  $\theta$ .
- \* Requires three forward passes on each step<sup>5</sup>.

---

<sup>5</sup>how many backwards?

**Table 1:** DQN targets

DQN	target $y$
Classic Deep Q-learning	$r + Q(s', \underset{a'}{\operatorname{argmax}} Q(s', a', \theta), \theta)$
With target-network	$r + Q^{\operatorname{target}}(s', \underset{a'}{\operatorname{argmax}} Q^{\operatorname{target}}(s', a', \theta^-), \theta^-)$
Double Deep Q-learning	$r + Q^{\operatorname{target}}(s', \underset{a'}{\operatorname{argmax}} Q(s', a', \theta), \theta^-)$

# Dueling DQN: Motivation

Note:

- \* In most states our choice of action does not affect future value.
- \* After finding  $Q(s, a)$  Q-learning still gains no information about  $Q(s, a')$  for  $a' \neq a$ .

# Dueling DQN: Motivation

Note:

- \* In most states our choice of action does not affect future value.
- \* After finding  $Q(s, a)$  Q-learning still gains no information about  $Q(s, a')$  for  $a' \neq a$ .

⇒ after trying an action in a bad state, Q-learning wants to try all other actions in this state.

# Dueling DQN: Motivation

Note:

- \* In most states our choice of action does not affect future value.
- \* After finding  $Q(s, a)$  Q-learning still gains no information about  $Q(s, a')$  for  $a' \neq a$ .

⇒ after trying an action in a bad state, Q-learning wants to try all other actions in this state.



Learning  $Q(s, a)$  should lead to learning  $V(s)$

# Advantage function

Define *advantage* function:

$$A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s)$$



# Advantage function

Define *advantage* function:

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

Note:

$$\begin{aligned}\mathbb{E}_{a \sim \pi} A^\pi(s, a) &= \mathbb{E}_{a \sim \pi} Q^\pi(s, a) - V^\pi(s) = \\ &= \mathbb{E}_{a \sim \pi} Q^\pi(s, a) - \mathbb{E}_{a \sim \pi} Q^\pi(s, a) = 0\end{aligned}$$

# Advantage function

Define *advantage* function:

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

Note:


$$\begin{aligned}\mathbb{E}_{a \sim \pi} A^\pi(s, a) &= \mathbb{E}_{a \sim \pi} Q^\pi(s, a) - V^\pi(s) = \\ &= \mathbb{E}_{a \sim \pi} Q^\pi(s, a) - \mathbb{E}_{a \sim \pi} Q^\pi(s, a) = 0\end{aligned}$$

Rewrite  $Q$ -function in terms of value of state:

$$Q^\pi(s, a) = V^\pi(s) + A^\pi(s, a)$$

# Dueling DQN (2016)

## Dueling Network Architectures for Deep Reinforcement Learning (2016)



Dueling Q-network architecture

## Struggling with identifiability

**Problem:**  $A(s, a)$  is not arbitrary. Recall  $\mathbb{E}_{a \sim \pi} A^\pi(s, a) = 0$ .

# Struggling with identifiability

**Problem:**  $A(s, a)$  is not arbitrary. Recall  $\mathbb{E}_{a \sim \pi} A^\pi(s, a) = 0$ .

For deterministic  $\pi(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$  this is equivalent to

$$\max_a A(s, a) = 0$$

# Struggling with identifiability

**Problem:**  $A(s, a)$  is not arbitrary. Recall  $\mathbb{E}_{a \sim \pi} A^\pi(s, a) = 0$ .

For deterministic  $\pi(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$  this is equivalent to

$$\max_a A(s, a) = 0$$

**Proposition:**

$$Q(s, a) = V(s) + A(s, a) - \max_a A(s, a)$$

# Struggling with identifiability

**Problem:**  $A(s, a)$  is not arbitrary. Recall  $\mathbb{E}_{a \sim \pi} A^\pi(s, a) = 0$ .

For deterministic  $\pi(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$  this is equivalent to

$$\max_a A(s, a) = 0$$

**Proposition:**

$$Q(s, a) = V(s) + A(s, a) - \max_a A(s, a)$$



$$Q(s, a) = V(s) + A(s, a) - \underset{a}{\operatorname{mean}} A(s, a)$$

suddenly worked better.

- ✓ Learning  $Q(s, a)$  leads to correcting  $V(s)$ .
- \* Only network architecture is changed.
- \* Double DQN still works for dueling architecture.



## Prioritized replay memory (2015)

---

In standard DQN we sample batch of transitions from replay memory uniformly.

- × Some transitions are more important than others
- × Replay memory is full of almost useless transitions

In standard DQN we sample batch of transitions from replay memory uniformly.

- × Some transitions are more important than others
- × Replay memory is full of almost useless transitions



$\delta = |y(s', r, done) - Q(s, a)|$  is  
a good proxy of transition importance

# Prioritized Replay Memory (2015)

Prioritized Experience Replay (2015):

$$p(\mathcal{T}) \propto \delta(\mathcal{T})^\alpha$$

Authors found  $\alpha \approx 0.6$  is a good universal value.

# Prioritized Replay Memory (2015)

Prioritized Experience Replay (2015):

$$p(\mathcal{T}) \propto \delta(\mathcal{T})^\alpha$$

Authors found  $\alpha \approx 0.6$  is a good universal value.

## Problems:

- × On each step this probability changes for all the replay memory <sup>6</sup>

---

<sup>6</sup>which capacity is on the order of 1M transitions

# Prioritized Replay Memory (2015)

Prioritized Experience Replay (2015):

$$p(\mathcal{T}) \propto \delta(\mathcal{T})^\alpha$$

Authors found  $\alpha \approx 0.6$  is a good universal value.

## Problems:

- × On each step this probability changes for all the replay memory <sup>6</sup>  
     $\approx$  on each step update  $\delta$  only for the sampled batch used for learning

---

<sup>6</sup>which capacity is on the order of 1M transitions

# Prioritized Replay Memory (2015)

Prioritized Experience Replay (2015):

$$p(\mathcal{T}) \propto \delta(\mathcal{T})^\alpha$$

Authors found  $\alpha \approx 0.6$  is a good universal value.

## Problems:

- × On each step this probability changes for all the replay memory <sup>6</sup>  
     $\approx$  on each step update  $\delta$  only for the sampled batch used for learning
- × Introduces **bias** (transitions are now sampled from hell knows what distribution).

---

<sup>6</sup>which capacity is on the order of 1M transitions

## Background: Importance Sampling

For arbitrary distribution  $q(x)$ :

$$\begin{aligned}\mathbb{E}_{p(x)} f(x) &= \int p(x) f(x) dx = \int \frac{q(x)}{q(x)} p(x) f(x) dx = \\ &= \int q(x) \frac{p(x)}{q(x)} f(x) dx = \mathbb{E}_{q(x)} \frac{p(x)}{q(x)} f(x)\end{aligned}$$



## Background: Importance Sampling

For arbitrary distribution  $q(x)$ :

$$\begin{aligned}\mathbb{E}_{p(x)} f(x) &= \int p(x) f(x) dx = \int \frac{q(x)}{q(x)} p(x) f(x) dx = \\ &= \int q(x) \frac{p(x)}{q(x)} f(x) dx = \mathbb{E}_{q(x)} \frac{p(x)}{q(x)} f(x)\end{aligned}$$

That's exactly what we want: to substitute expectation of loss over uniform sampling from experience replay to expectation over our own prioritized distribution!

## Background: Importance Sampling

For arbitrary distribution  $q(x)$ :

$$\begin{aligned}\mathbb{E}_{p(x)} f(x) &= \int p(x) f(x) dx = \int \frac{q(x)}{q(x)} p(x) f(x) dx = \\ &= \int q(x) \frac{p(x)}{q(x)} f(x) dx = \mathbb{E}_{q(x)} \frac{p(x)}{q(x)} f(x)\end{aligned}$$

That's exactly what we want: to substitute expectation of loss ( $f(x)$ ) over uniform sampling from experience replay ( $p(x)$ ) to expectation over our own prioritized distribution ( $q(x)$ ) !

# Applying Importance Sampling

If  $N$  is replay memory capacity:

$$L = \mathbb{E}_{\mathcal{T} \sim \text{uniform}} (y - Q(s, a))^2 = \mathbb{E}_{\mathcal{T} \sim \text{prioritized}} \frac{1}{Np(\mathcal{T})} (y - Q(s, a))^2$$

IS just adds weights to our batch:

$$w_i = \frac{1}{Np(\mathcal{T}_i)}$$

**Problem:** at the beginning these weights might not be that relevant, yet slowing down learning.

**Problem:** at the beginning these weights might not be that relevant, yet slowing down learning.

Let's smooth them at the beginning of learning:

$$L = \mathbb{E}_{\mathcal{T} \sim \text{prioritized}} \left( \frac{1}{Np(\mathcal{T})} \right)^{\beta} (y - Q(s, a))^2,$$

where  $\beta$  changes from 0.4 to 1 linearly during first 100 000 frames.

- \* Weights significantly vary scale of loss function. Constant learning rate might be inappropriate.

*Hint:*<sup>7</sup> normalize weights by dividing on  $\max_i w_i$ .

---

<sup>7</sup>max taken over all replay memory. Yet in some implementations it is taken over current batch

- \* Weights significantly vary scale of loss function. Constant learning rate might be inappropriate.

*Hint:*<sup>7</sup> normalize weights by dividing on  $\max_i w_i$ .

- \*  $\min(1, |\delta|)$  is used instead of  $|\delta|$  for stabilization purposes.

---

<sup>7</sup>max taken over all replay memory. Yet in some implementations it is taken over current batch

- \* Weights significantly vary scale of loss function. Constant learning rate might be inappropriate.

*Hint:*<sup>7</sup> normalize weights by dividing on  $\max_i w_i$ .

- \*  $\min(1, |\delta|)$  is used instead of  $|\delta|$  for stabilization purposes.
- \* new transitions are stored with maximum priority.

---

<sup>7</sup>max taken over all replay memory. Yet in some implementations it is taken over current batch



## Noisy networks for exploration (2017)

---

**Problem:**  $\epsilon$ -greedy exploration is *state-independent*.

**Problem:**  $\epsilon$ -greedy exploration is *state-independent*.



Add parametric noise to the weights of Q-network

# Noisy Nets (2017)

**Problem:**  $\epsilon$ -greedy exploration is *state-independent*.



Add parametric noise to the weights of Q-network

Noisy Nets for Exploration (2017):

$$w_i = \mu_i + \sigma_i * \varepsilon_i, \quad \varepsilon \sim \mathcal{N}(0, 1)$$

- \*  $\mu_i, \sigma_i$  are both learnable parameters.
- \* all weights are independent random variables
- \* use policy  $\pi(s) = \underset{a}{\operatorname{argmax}} Q(s, a, \mu, \sigma, \varepsilon)$

Formally, our loss<sup>8</sup> is now:

$$\mathbb{E}_{\varepsilon} \mathbb{E}_{\mathcal{T}} (Q(s, a, \theta, \varepsilon) - y(\mathcal{T}))^2$$

---

<sup>8</sup>Noisy Net is not a bayesian NN as it does not model probability; loss minimization is also not an upper bound optimization

Formally, our loss<sup>8</sup> is now:

$$\mathbb{E}_{\varepsilon} \mathbb{E}_{\mathcal{T}} (Q(s, a, \theta, \varepsilon) - y(\mathcal{T}))^2$$

**Problem:**  $y$  also depends on stochastic  $Q$ -function.

---

<sup>8</sup>Noisy Net is not a bayesian NN as it does not model probability; loss minimization is also not an upper bound optimization

Formally, our loss<sup>8</sup> is now:

$$\mathbb{E}_{\varepsilon} \mathbb{E}_{\mathcal{T}} (Q(s, a, \theta, \varepsilon) - y(\mathcal{T}))^2$$

**Problem:**  $y$  also depends on stochastic  $Q$ -function.

\* use different noise samples for it:

$$y = r + Q(s', \underset{a'}{\operatorname{argmax}} Q(s', a', \varepsilon''), \varepsilon')$$

---

<sup>8</sup>Noisy Net is not a bayesian NN as it does not model probability; loss minimization is also not an upper bound optimization

# Noise generation optimization

**Problem:** noise generation turns to be a bottleneck in terms of wall-clock time.

- $MN + N$  samples required for linear layer mapping  $M$  features to  $N$ .



# Noise generation optimization

**Problem:** noise generation turns to be a bottleneck in terms of wall-clock time.

- $MN + N$  samples required for linear layer mapping  $M$  features to  $N$ .

**Alternative:**

- generate  $M$  noises  $\varepsilon_m$  and  $N$  noises  $\varepsilon_n$
- consider weight noise  $\varepsilon_{mn} = f(\varepsilon_m)f(\varepsilon_n)$ , where  $f$  is scaling function (signed square root)

# Noise generation optimization

**Problem:** noise generation turns to be a bottleneck in terms of wall-clock time.

- $MN + N$  samples required for linear layer mapping  $M$  features to  $N$ .

**Alternative:**

- generate  $M$  noises  $\varepsilon_m$  and  $N$  noises  $\varepsilon_n$
- consider weight noise  $\varepsilon_{mn} = f(\varepsilon_m)f(\varepsilon_n)$ , where  $f$  is scaling function (signed square root)
- $N$  more noises for bias <sup>9</sup>

---

<sup>9</sup>authors also scale them with  $f$

# Noise generation optimization

**Problem:** noise generation turns to be a bottleneck in terms of wall-clock time.

- $MN + N$  samples required for linear layer mapping  $M$  features to  $N$ .

**Alternative:**

- generate  $M$  noises  $\varepsilon_m$  and  $N$  noises  $\varepsilon_n$
  - consider weight noise  $\varepsilon_{mn} = f(\varepsilon_m)f(\varepsilon_n)$ , where  $f$  is scaling function (signed square root)
  - $N$  more noises for bias <sup>9</sup>
- ✓ just  $M + 2N$  noise samples are needed.

---

<sup>9</sup>authors also scale them with  $f$

# Noise generation optimization

**Problem:** noise generation turns to be a bottleneck in terms of wall-clock time.

- $MN + N$  samples required for linear layer mapping  $M$  features to  $N$ .

**Alternative:**

- generate  $M$  noises  $\varepsilon_m$  and  $N$  noises  $\varepsilon_n$
  - consider weight noise  $\varepsilon_{mn} = f(\varepsilon_m)f(\varepsilon_n)$ , where  $f$  is scaling function (signed square root)
  - $N$  more noises for bias <sup>9</sup>
- ✓ just  $M + 2N$  noise samples are needed.
- \* for whole batch!<sup>10</sup>

---

<sup>9</sup>authors also scale them with  $f$

<sup>10</sup>is this theoretically coherent?

- ✓ No hyperparameters
  - \* Except where to put noise in the network... Convolution layers better leave deterministic<sup>11</sup>.

---

<sup>11</sup>why?

- ✓ No hyperparameters
  - \* Except where to put noise in the network... Convolution layers better leave deterministic<sup>11</sup>.
- ✓ noise magnitude  $\sigma$  will (hopefully<sup>12</sup>) vanish state-dependently through the learning process

---

<sup>11</sup>why?

<sup>12</sup>on practice, behaves very differently from game to game

- ✓ No hyperparameters
  - \* Except where to put noise in the network... Convolution layers better leave deterministic<sup>11</sup>.
- ✓ noise magnitude  $\sigma$  will (hopefully<sup>12</sup>) vanish state-dependently through the learning process
  - \* yet  $w_i = \mu_i$  can be used for exploitation purposes

---

<sup>11</sup>why?

<sup>12</sup>on practice, behaves very differently from game to game

- ✓ No hyperparameters
  - \* Except where to put noise in the network... Convolution layers better leave deterministic<sup>11</sup>.
- ✓ noise magnitude  $\sigma$  will (hopefully<sup>12</sup>) vanish state-dependently through the learning process
  - \* yet  $w_i = \mu_i$  can be used for exploitation purposes
- ✓ almost random behavior at the beginning

---

<sup>11</sup>why?

<sup>12</sup>on practice, behaves very differently from game to game



# Application tips

- ✓ No hyperparameters
  - \* Except where to put noise in the network... Convolution layers better leave deterministic<sup>11</sup>.
- ✓ noise magnitude  $\sigma$  will (hopefully<sup>12</sup>) vanish state-dependently through the learning process
  - \* yet  $w_i = \mu_i$  can be used for exploitation purposes
- ✓ almost random behavior at the beginning
  - \* yet  $\epsilon$ -greedy strategy may also be used

---

<sup>11</sup>why?

<sup>12</sup>on practice, behaves very differently from game to game

## Categorical DQN (2017)

---

# Motivation

Consider a state where you get 1000 or -1000 with probabilities 0.5.

# Motivation

Consider a state where you get 1000 or -1000 with probabilities 0.5.  
Q-network would say value of state is 0.

# Motivation

Consider a state where you get 1000 or -1000 with probabilities 0.5.  
Q-network would say value of state is 0. **But you never really get 0.**

# Motivation

Consider a state where you get 1000 or -1000 with probabilities 0.5.  
Q-network would say value of state is 0. **But you never really get 0.**



Learn a **distribution** over future reward instead of it's expectation.

Recall

$$Q^{\pi}(s, a) = \mathbb{E} \sum_t r(s_t) \mid s, a$$

# Value Distribution

Recall

$$Q^{\pi}(s, a) = \mathbb{E} \sum_t r(s_t) \mid s, a$$

A Distributional Perspective on Reinforcement Learning (2017):

For fixed policy  $\pi$  let's define *value distribution*:

## Value distribution

Let's define value distribution as distribution of

$$Z^{\pi}(s, a) = \sum_t r(s_t) \mid s, a$$



# Value Distribution

Recall

$$Q^\pi(s, a) = \mathbb{E} \sum_t r(s_t) \mid s, a$$

A Distributional Perspective on Reinforcement Learning (2017):

For fixed policy  $\pi$  let's define *value distribution*:

## Value distribution

Let's define value distribution as distribution of

$$Z^\pi(s, a) = \sum_t r(s_t) \mid s, a$$

**! It's a random variable!**

Value distribution satisfies a *recursive distributional equation*:

$$Z^\pi(s, a) \stackrel{\text{D}}{=} r(s, a) + \gamma Z^\pi(s', \pi(s'))$$

# Dynamic programming for value distribution

Value distribution satisfies a *recursive distributional equation*:

$$Z^\pi(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^\pi(s', \pi(s'))$$

Just equivalence of c.d.f. of left and right part :}

# Dynamic programming for value distribution

Value distribution satisfies a *recursive distributional equation*:

$$Z^\pi(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^\pi(s', \pi(s'))$$

Just equivalence of c.d.f. of left and right part :}

**Question:** will point iteration be a contraction mapping for some metric in the space of value distributions?

# Dynamic programming for value distribution

Value distribution satisfies a *recursive distributional equation*:

$$Z^\pi(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^\pi(s', \pi(s'))$$

Just equivalence of c.d.f. of left and right part :}

**Question:** will point iteration be a contraction mapping for some metric in the **space of value distributions**?

# Dynamic programming for value distribution

Value distribution satisfies a *recursive distributional equation*:

$$Z^\pi(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^\pi(s', \pi(s'))$$

Just equivalence of c.d.f. of left and right part :}

**Question:** will point iteration be a contraction mapping for some metric in the space of value distributions?

- ✓ yes, for  $d(Z_1, Z_2) = \sup_{s, a} \mathcal{W}(Z_1(s, a), Z_2(s, a))$ , where  $\mathcal{W}$  is Wasserstein distance between two random variables.

Analogically:  $\pi^*(s) = \max_a \mathbb{E} Z^{\pi^*}(s, a)$

# Distributional Bellman Equation

Analogically:  $\pi^*(s) = \max_a \mathbb{E} Z^{\pi^*}(s, a)$

## Distributional Bellman equation

$$Z^{\pi^*}(s, a) \stackrel{\text{D}}{=} r(s, a) + \gamma Z^{\pi^*}(s', \pi^*(s'))$$



# Distributional Bellman Equation

Analogically:  $\pi^*(s) = \max_a \mathbb{E} Z^{\pi^*}(s, a)$

## Distributional Bellman equation

$$Z^{\pi^*}(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^{\pi^*}(s', \max_{a'} \mathbb{E} Z^{\pi^*}(s', a'))$$

# Distributional Bellman Equation

Analogically:  $\pi^*(s) = \max_a \mathbb{E} Z^{\pi^*}(s, a)$

## Distributional Bellman equation

$$Z^{\pi^*}(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^{\pi^*}(s', \max_{a'} \mathbb{E} Z^{\pi^*}(s', a'))$$

**Question:** will point iteration be a contraction mapping for some metric in the space of value distributions? <sup>13</sup>

---

<sup>13</sup>and why are we asking this again?

# Distributional Bellman Equation

Analogically:  $\pi^*(s) = \max_a \mathbb{E} Z^{\pi^*}(s, a)$

## Distributional Bellman equation

$$Z^{\pi^*}(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^{\pi^*}(s', \max_{a'} \mathbb{E} Z^{\pi^*}(s', a'))$$

**Question:** will point iteration be a contraction mapping for some metric in the space of value distributions? <sup>13</sup>

× no, it will not.

---

<sup>13</sup>and why are we asking this again?

# Distributional Bellman Equation

Analogically:  $\pi^*(s) = \max_a \mathbb{E} Z^{\pi^*}(s, a)$

## Distributional Bellman equation

$$Z^{\pi^*}(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^{\pi^*}(s', \max_{a'} \mathbb{E} Z^{\pi^*}(s', a'))$$

**Question:** will point iteration be a contraction mapping for some metric in the space of value distributions? <sup>13</sup>

- × no, it will not.
- × there may be no fixed point at all

---

<sup>13</sup>and why are we asking this again?

# Distributional Bellman Equation


Analogically:  $\pi^*(s) = \max_a \mathbb{E} Z^{\pi^*}(s, a)$

## Distributional Bellman equation

$$Z^{\pi^*}(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^{\pi^*}(s', \max_{a'} \mathbb{E} Z^{\pi^*}(s', a'))$$

**Question:** will point iteration be a contraction mapping for some metric in the space of value distributions? <sup>13</sup>

- × no, it will not.
- × there may be no fixed point at all
- × and existence of one doesn't guarantee convergence to it



---

<sup>13</sup>and why are we asking this again?

# We'll do it anyway!

Let's do point iteration anyway! Our wish:

$$p(Z_{t+1}(s, a)) \leftarrow p\left(r(s, a) + \gamma Z_t\left[s', \max_{a'} \mathbb{E} Z_t(s', a')\right]\right)$$

# We'll do it anyway!

Let's do point iteration anyway! Our wish:

$$p(Z_{t+1}(s, a)) \leftarrow p\left(r(s, a) + \gamma Z_t \left[s', \max_{a'} \mathbb{E} Z_t(s', a')\right]\right)$$

**Problems:**

?  $p(Z_t)$  is some distribution on  $\mathbb{R}$ . How do we represent it?

# We'll do it anyway!

Let's do point iteration anyway! Our wish:

$$p(Z_{t+1}(s, a)) \leftarrow p\left(r(s, a) + \gamma Z_t\left[s', \max_{a'} \mathbb{E} Z_t(s', a')\right]\right)$$

**Problems:**

?  $p(Z_t)$  is some distribution on  $\mathbb{R}$ . How do we represent it?

✓ use some distribution family as approximation:  $p(Z_t(s, a)) \approx Z_\theta$



# We'll do it anyway!

Let's do point iteration anyway! Our wish:

$$p(Z_{t+1}(s, a)) \leftarrow p\left(r(s, a) + \gamma Z_t\left[s', \max_{a'} \mathbb{E} Z_t(s', a')\right]\right)$$

## Problems:

?  $p(Z_t)$  is some distribution on  $\mathbb{R}$ . How do we represent it?

✓ use some distribution family as approximation:  $p(Z_t(s, a)) \approx Z_\theta$

? we have only samples of  $s'$  and  $r(s, a)$

# We'll do it anyway!

Let's do point iteration anyway! Our wish:

$$p(Z_{t+1}(s, a)) \leftarrow p\left(r(s, a) + \gamma Z_t\left[s', \max_{a'} \mathbb{E} Z_t(s', a')\right]\right)$$

## Problems:

?  $p(Z_t)$  is some distribution on  $\mathbb{R}$ . How do we represent it?

✓ use some distribution family as approximation:  $p(Z_t(s, a)) \approx Z_\theta$

? we have only samples of  $s'$  and  $r(s, a)$

\* in DQN case we optimized  $L = \mathbb{E}_{\mathcal{T}}(y - Q_\theta)^2$

# We'll do it anyway!

Let's do point iteration anyway! Our wish:

$$p(Z_{t+1}(s, a)) \leftarrow p\left(r(s, a) + \gamma Z_t\left[s', \max_{a'} \mathbb{E} Z_t(s', a')\right]\right)$$

## Problems:

?  $p(Z_t)$  is some distribution on  $\mathbb{R}$ . How do we represent it?

✓ use some distribution family as approximation:  $p(Z_t(s, a)) \approx Z_\theta$

? we have only samples of  $s'$  and  $r(s, a)$

\* in DQN case we optimized  $L = \mathbb{E}_{\mathcal{T}}(y - Q_\theta)^2$

✓ let's optimize  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta)$ , where  $\mathcal{D}$  is some divergence

# We'll do it anyway!

Let's do point iteration anyway! Our wish:

$$p(Z_{t+1}(s, a)) \leftarrow p\left(r(s, a) + \gamma Z_t\left[s', \max_{a'} \mathbb{E} Z_t(s', a')\right]\right)$$

## Problems:

?  $p(Z_t)$  is some distribution on  $\mathbb{R}$ . How do we represent it?

✓ use some distribution family as approximation:  $p(Z_t(s, a)) \approx Z_\theta$

? we have only samples of  $s'$  and  $r(s, a)$

\* in DQN case we optimized  $L = \mathbb{E}_{\mathcal{T}}(y - Q_\theta)^2$

✓ let's optimize  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta)$ , where  $\mathcal{D}$  is some divergence

$Z_{t+1} \stackrel{D}{=} r(s, a) + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$  is a convolution involving MDP transition probability  $p(s' \mid s, a)$ :

$$p(Z_{t+1}) = \sum_{s'} p(y \mid s') p(s' \mid s, a)$$

where  $y(s') = r(s') + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$ .

$Z_{t+1} \stackrel{D}{=} r(s, a) + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$  is a convolution involving MDP transition probability  $p(s' \mid s, a)$ :

$$p(Z_{t+1}) = \sum_{s'} p(y \mid s') p(s' \mid s, a)$$

where  $y(s') = r(s') + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$ .

**Note:** for fixed  $Z_t(s, a)$  and given  $s'$  function  $y(s')$  is deterministic!

$Z_{t+1} \stackrel{D}{=} r(s, a) + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$  is a convolution involving MDP transition probability  $p(s' \mid s, a)$ :

$$p(Z_{t+1}) = \sum_{s'} p(y \mid s') p(s' \mid s, a)$$

where  $y(s') = r(s') + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$ .

**Note:** for fixed  $Z_t(s, a)$  and given  $s'$  function  $y(s')$  is deterministic!

$\Rightarrow$  for given  $p(Z_t(s, a))$  we can get  $p(y \mid s')$

$Z_{t+1} \stackrel{D}{=} r(s, a) + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$  is a convolution involving MDP transition probability  $p(s' \mid s, a)$ :

$$p(Z_{t+1}) = \sum_{s'} p(y \mid s') p(s' \mid s, a)$$

where  $y(s') = r(s') + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$ .

**Note:** for fixed  $Z_t(s, a)$  and given  $s'$  function  $y(s')$  is deterministic!

$\Rightarrow$  for given  $p(Z_t(s, a))$  we can get  $p(y \mid s')$

**Problem:** and what?



$Z_{t+1} \stackrel{D}{=} r(s, a) + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$  is a convolution involving MDP transition probability  $p(s' | s, a)$ :

$$p(Z_{t+1}) = \sum_{s'} p(y | s') p(s' | s, a) = \mathbb{E}_{\mathcal{T}} p(y | s')$$

where  $y(s') = r(s') + \gamma Z_t(s', \max_a \mathbb{E} Z_t(s, a))$ .

**Note:** for fixed  $Z_t(s, a)$  and given  $s'$  function  $y(s')$  is deterministic!

$\Rightarrow$  for given  $p(Z_t(s, a))$  we can get  $p(y | s')$

**Problem:** and what?

Our loss is  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta) = \mathcal{D}(\mathbb{E}_{\mathcal{T}} p(y \mid s') \parallel Z_\theta)$ .

Our loss is  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta) = \mathcal{D}(\mathbb{E}_{\mathcal{T}} p(y \mid s') \parallel Z_\theta)$ .

\* we are free to choose  $\mathcal{D}$ !

Our loss is  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta) = \mathcal{D}(\mathbb{E}_{\mathcal{T}} p(y \mid s') \parallel Z_\theta)$ .

\* we are free to choose  $\mathcal{D}$ !

$$\text{KL}(p(Z_{t+1}) \parallel Z_\theta) = - \int p(Z_{t+1}) \log Z_\theta + C = - \int \mathbb{E}_{\mathcal{T}} p(y \mid s') \log Z_\theta + C =$$

Our loss is  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta) = \mathcal{D}(\mathbb{E}_{\mathcal{T}} p(y \mid s') \parallel Z_\theta)$ .

\* we are free to choose  $\mathcal{D}$ !

$$\begin{aligned}\text{KL}(p(Z_{t+1}) \parallel Z_\theta) &= - \int p(Z_{t+1}) \log Z_\theta + C = - \int \mathbb{E}_{\mathcal{T}} p(y \mid s') \log Z_\theta + C = \\ &= - \mathbb{E}_{\mathcal{T}} \int p(y \mid s') \log Z_\theta + C = \mathbb{E}_{\mathcal{T}} \text{KL}(p(y \mid s') \parallel Z_\theta) + C\end{aligned}$$

Our loss is  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta) = \mathcal{D}(\mathbb{E}_{\mathcal{T}} p(y \mid s') \parallel Z_\theta)$ .

\* we are free to choose  $\mathcal{D}$ !

$$\begin{aligned}\text{KL}(p(Z_{t+1}) \parallel Z_\theta) &= - \int p(Z_{t+1}) \log Z_\theta + C = - \int \mathbb{E}_{\mathcal{T}} p(y \mid s') \log Z_\theta + C = \\ &= - \mathbb{E}_{\mathcal{T}} \int p(y \mid s') \log Z_\theta + C = \mathbb{E}_{\mathcal{T}} \text{KL}(p(y \mid s') \parallel Z_\theta) + C\end{aligned}$$

✓ this can be evaluated through Monte-Carlo!

Our loss is  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta) = \mathcal{D}(\mathbb{E}_{\mathcal{T}} p(y \mid s') \parallel Z_\theta)$ .

\* we are free to choose  $\mathcal{D}$ !

$$\begin{aligned}\text{KL}(p(Z_{t+1}) \parallel Z_\theta) &= - \int p(Z_{t+1}) \log Z_\theta + C = - \int \mathbb{E}_{\mathcal{T}} p(y \mid s') \log Z_\theta + C = \\ &= - \mathbb{E}_{\mathcal{T}} \int p(y \mid s') \log Z_\theta + C = \mathbb{E}_{\mathcal{T}} \text{KL}(p(y \mid s') \parallel Z_\theta) + C\end{aligned}$$

✓ this can be evaluated through Monte-Carlo!

× trick doesn't work for other divergences!

Our loss is  $L = \mathcal{D}(p(Z_{t+1}) \parallel Z_\theta) = \mathcal{D}(\mathbb{E}_{\mathcal{T}} p(y \mid s') \parallel Z_\theta)$ .

\* we are free to choose  $\mathcal{D}$ !

$$\begin{aligned}\text{KL}(p(Z_{t+1}) \parallel Z_\theta) &= - \int p(Z_{t+1}) \log Z_\theta + C = - \int \mathbb{E}_{\mathcal{T}} p(y \mid s') \log Z_\theta + C = \\ &= - \mathbb{E}_{\mathcal{T}} \int p(y \mid s') \log Z_\theta + C = \mathbb{E}_{\mathcal{T}} \text{KL}(p(y \mid s') \parallel Z_\theta) + C\end{aligned}$$

✓ this can be evaluated through Monte-Carlo!

× trick doesn't work for other divergences!

\* KL requires  $Z_{t+1}$  and  $Z_\theta$  share domain.



# Distribution family

Options:

- Gaussian mixture
- Discrete

# Distribution family

Options:

- Gaussian mixture
- Discrete ✓
  - \* KL-divergence of two discrete distributions is a simple cross-entropy

Options:

- Gaussian mixture
- Discrete ✓
  - \* KL-divergence of two discrete distributions is a simple cross-entropy

Let's  $\mathcal{P}$  be a family of categorical distribution on the grid from  $V_{\min}$  to  $V_{\max}$  with  $N$  atoms (outcomes).

# Distribution family

Options:

- Gaussian mixture
- Discrete ✓
  - \* KL-divergence of two discrete distributions is a simple cross-entropy

Let's  $\mathcal{P}$  be a family of categorical distribution on the grid from  $V_{\min}$  to  $V_{\max}$  with  $N$  atoms (outcomes).

## Parametrization:

For each action our neural network  $Z(s, a)$  outputs  $N$  numbers, summing into 1

## Calculating target

Suppose you have transition  $(s, a, r, s', done)$ ,  $Z(s, a) \in \mathcal{P}$ . Then:

$$y(s') = r + \gamma Z(s', \max_{a'} \mathbb{E} Z(s', a'))$$

## Calculating target





Suppose you have transition  $(s, a, r, s', done)$ ,  $Z(s, a) \in \mathcal{P}$ . Then:

$$y(s') = r + \gamma Z(s', \max_{a'} \mathbb{E} Z(s', a')) \in \mathcal{P}?$$

# Calculating target

Suppose you have transition  $(s, a, r, s', done)$ ,  $Z(s, a) \in \mathcal{P}$ . Then:

$$y(s') = \Pr \left[ r + \gamma Z(s', \max_{a'} \mathbb{E} Z(s', a')) \right] \in \mathcal{P}$$



# How it looks like

Failed to insert video into beamer ;o)



## Rainbow DQN (2018)

---

Blend them all!



Multistep  
DQN

DQN



Double  
DQN



Prioritized  
Replay



Noisy  
Net



Categorical  
DQN

Dueling  
DQN



# Multistep DQN: Motivation

Recall our target in classic DQN:

$$y = r + \gamma \max_{a'} Q(s', a')$$

If we have nonzero reward at the end of  $M$ -step game, we need at least  $M$  iterations of Q-learning to «propagate» this reward to all visited states.

# Multistep DQN: Motivation

Recall our target in classic DQN:

$$y = r + \gamma \max_{a'} Q(s', a')$$

If we have nonzero reward at the end of  $M$ -step game, we need at least  $M$  iterations of Q-learning to «propagate» this reward to all visited states.



Look more than one step ahead!

# Multistep DQN: Realisation

- work with transitions  $(s, a, r, r', r'', \dots, r^{(M-1)}, s^{(M)}, done)$

# Multistep DQN: Realisation

- work with transitions  $(s, a, r, r', r'', \dots, r^{(M-1)}, s^{(M)}, done)$
- transition reward  $R = r + \gamma r' + \gamma^2 r'' + \dots + \gamma^{M-1} r^{M-1}$

# Multistep DQN: Realisation

- work with transitions  $(s, a, r, r', r'', \dots, r^{(M-1)}, s^{(M)}, done)$
- transition reward  $R = r + \gamma r' + \gamma^2 r'' + \dots + \gamma^{M-1} r^{M-1}$
- use new target:

$$y = R + \gamma^M \max_{a^{(M)}} Q(s^{(M)}, a^{(M)})$$

# Multistep DQN: Realisation

- work with transitions  $(s, a, r, r', r'', \dots, r^{(M-1)}, s^{(M)}, done)$
- transition reward  $R = r + \gamma r' + \gamma^2 r'' + \dots + \gamma^{M-1} r^{M-1}$
- use new target:

$$y = R + \gamma^M \max_{a^{(M)}} Q(s^{(M)}, a^{(M)})$$

× formally can be used only with on-policy algorithms <sup>14</sup>

---

<sup>14</sup>why?



# Multistep DQN: Realisation

- work with transitions  $(s, a, r, r', r'', \dots, r^{(M-1)}, s^{(M)}, done)$
- transition reward  $R = r + \gamma r' + \gamma^2 r'' + \dots + \gamma^{M-1} r^{M-1}$
- use new target:

$$y = R + \gamma^M \max_{a^{(M)}} Q(s^{(M)}, a^{(M)})$$

- × formally can be used only with on-policy algorithms <sup>14</sup>
- × the further we look the worse  $y$  approximates  $Q^{\pi^*}(s, a)$   
⇒ number of steps should be chosen carefully.

---

<sup>14</sup>why?

# Multistep Categorical DQN

Recall categorical DQN target:

$$y = \Pr \left[ r + \gamma Z(s', \underset{a'}{\operatorname{argmax}} \mathbb{E} Z(s', a')) \right]$$

# Multistep Categorical DQN

Recall categorical DQN target:

$$y = \Pr \left[ r + \gamma Z(s', \underset{a'}{\operatorname{argmax}} \mathbb{E} Z(s', a')) \right]$$

It can also be made multi-step:

$$y = \Pr \left[ R + \gamma^M Z(s^{(M)}, \underset{a^{(M)}}{\operatorname{argmax}} \mathbb{E} Z(s^{(M)}, a^{(M)})) \right]$$

# Multistep Categorical DQN

Recall categorical DQN target:

$$y = \Pr \left[ r + \gamma Z(s', \underset{a'}{\operatorname{argmax}} \mathbb{E} Z(s', a')) \right]$$

It can also be made multi-step:

$$y = \Pr \left[ R + \gamma^M Z(s^{(M)}, \underset{a^{(M)}}{\operatorname{argmax}} \mathbb{E} Z(s^{(M)}, a^{(M)})) \right]$$

Loss stays the same:

$$L = \text{KL}(p(y) \parallel p(Z))$$

# Dueling Categorical DQN

Recall dueling DQN:

$$Q(s, a) = V(s) + A(s, a) - \underset{a}{\text{mean}} A(s, a)$$

# Dueling Categorical DQN

Recall dueling DQN:

$$Q(s, a) = V(s) + A(s, a) - \underset{a}{\text{mean}} A(s, a)$$

Let's make our  $Z(s, a)$  (modeling categorical distribution with  $N$  atoms) in dueling way:

$$Z(s, a) = V_N(s) + A_N(s, a) - \underset{a}{\text{mean}} A_N(s, a)$$

where  $V_N(s)$  and  $A_N(s, a)$  are categorical  $N$ -atomized distributions.

# Dueling Categorical DQN

Recall dueling DQN:

$$Q(s, a) = V(s) + A(s, a) - \underset{a}{\text{mean}} A(s, a)$$

Let's make our  $Z(s, a)$  (modeling categorical distribution with  $N$  atoms) in dueling way:

$$Z(s, a) = \text{softmax}(V_N(s) + A_N(s, a) - \underset{a}{\text{mean}} A_N(s, a))$$

where  $V_N(s)$  and  $A_N(s, a)$  are arbitrary  $N$  numbers<sup>15</sup>.



---

<sup>15</sup>why couldn't we only add softmax?

# Double dueling multi-step noised categorical DQN with prioritized replay AKA Rainbow

Rainbow: Combining Improvements in Deep Reinforcement Learning (2018):

**Dueling + Multistep + Categorical + DQN +**



# Double dueling multi-step noised categorical DQN with prioritized replay AKA Rainbow

Rainbow: Combining Improvements in Deep Reinforcement Learning (2018):

**Dueling + Multistep + Categorical + DQN +**

- **Double:** use target network to evaluate

$$y = R + \gamma^N Z^{\text{target}}(s^{(N)}, \underset{a^{(N)}}{\operatorname{argmax}} Z(s^{(N)}, a^{(N)}))$$

# Double dueling multi-step noised categorical DQN with prioritized replay AKA Rainbow

Rainbow: Combining Improvements in Deep Reinforcement Learning (2018):

**Dueling + Multistep + Categorical + DQN +**

- **Double:** use target network to evaluate

$$y = R + \gamma^N Z^{\text{target}}(s^{(N)}, \underset{a^{(N)}}{\operatorname{argmax}} Z(s^{(N)}, a^{(N)}))$$

- **Noisy:** add noise to all fully connected layers

# Double dueling multi-step noised categorical DQN with prioritized replay AKA Rainbow

Rainbow: Combining Improvements in Deep Reinforcement Learning (2018):

**Dueling + Multistep + Categorical + DQN +**

- **Double:** use target network to evaluate


$$y = R + \gamma^N Z^{\text{target}}(s^{(N)}, \underset{a^{(N)}}{\operatorname{argmax}} Z(s^{(N)}, a^{(N)}))$$

- **Noisy:** add noise to all fully connected layers
- **Prioritized Replay:** just use it<sup>16</sup>

---

<sup>16</sup>guess proxy of transition priority

# Do we really need all this?



- \* all improvements are important as they address different problems

- \* all improvements are important as they address different problems
- × a lot of hyperparameters

- \* all improvements are important as they address different problems
- × a lot of hyperparameters
- ? Allegedly 10 hours for 7M frames on single GPU
  - :( I can't reproduce <sup>17</sup>

---

<sup>17</sup>10 hours for 3M. Noise generation seems to be a problem!

**NEXT: see pt.2 for Policy Gradient algorithms**