

Deep Reinforcement Learning

Overview of main articles

Part 2. Policy gradient algorithms

Sergey Ivanov

February 24, 2019

MSU

Basic policy gradient methods

REINFORCE

Baselines introduction

Actor-Critic

Generalized Advantage Estimation (GAE) (2018)

Trust Region Policy Optimization (TRPO) (2017)

Basic policy gradient methods

Recall RL goal:

$$\mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R \rightarrow \max_{\theta}$$

Direct optimization

Recall RL goal:

$$\mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R \rightarrow \max_{\theta}$$

Let's optimize our goal directly!

$$\nabla_{\theta} \mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R \quad \text{--- ?}$$

Direct optimization

Recall RL goal:

$$\mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R \rightarrow \max_{\theta}$$

Let's optimize our goal directly!

$$\nabla_{\theta} \mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R \quad \text{--- ?}$$

Direct optimization

Recall RL goal:

$$\mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R \rightarrow \max_{\theta}$$

Let's optimize our goal directly!

$$\nabla_{\theta} \mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R \quad \text{— ?}$$

Options:

- * Metaheuristics
- * Log-derivative trick¹.

¹aka REINFORCE

Stochastic estimators optimization via log-derivative trick

$$\nabla_{\theta} \mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \nabla_{\theta} \int \pi(x, \theta) f(x) dx$$

Stochastic estimators optimization via log-derivative trick

$$\nabla_{\theta} \mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \nabla_{\theta} \int \pi(x, \theta) f(x) dx = \left\{ \text{👤} \right\} = \int \nabla_{\theta} \pi(x, \theta) f(x) dx$$

Stochastic estimators optimization via log-derivative trick

$$\nabla_{\theta} \mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \nabla_{\theta} \int \pi(x, \theta) f(x) dx = \left\{ \text{👤} \right\} = \int \nabla_{\theta} \pi(x, \theta) f(x) dx$$

Problem: and what?

Stochastic estimators optimization via log-derivative trick

$$\nabla_{\theta} \mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \nabla_{\theta} \int \pi(x, \theta) f(x) dx = \left\{ \text{👤} \right\} = \int \nabla_{\theta} \pi(x, \theta) f(x) dx$$

Problem: and what?

Log-derivative trick

$$\nabla_{\theta} \pi(\theta) = \pi(\theta) \nabla_{\theta} \log \pi(\theta)$$

Stochastic estimators optimization via log-derivative trick

$$\nabla_{\theta} \mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \nabla_{\theta} \int \pi(x, \theta) f(x) dx = \left\{ \text{👤} \right\} = \int \nabla_{\theta} \pi(x, \theta) f(x) dx =$$

Problem: and what?

Log-derivative trick

$$\nabla_{\theta} \pi(\theta) = \pi(\theta) \nabla_{\theta} \log \pi(\theta)$$

$$= \int \pi(x, \theta) \nabla_{\theta} \log \pi(x, \theta) f(x) dx$$

Stochastic estimators optimization via log-derivative trick

$$\nabla_{\theta} \mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \nabla_{\theta} \int \pi(x, \theta) f(x) dx = \left\{ \text{👤} \right\} = \int \nabla_{\theta} \pi(x, \theta) f(x) dx =$$

Problem: and what?

Log-derivative trick

$$\nabla_{\theta} \pi(\theta) = \pi(\theta) \nabla_{\theta} \log \pi(\theta)$$

$$= \int \pi(x, \theta) \nabla_{\theta} \log \pi(x, \theta) f(x) dx = \mathbb{E}_{x \sim \pi(x, \theta)} \nabla_{\theta} \log \pi(x, \theta) f(x)$$

From importance sampling point of view

Recall Importance Sampling. For arbitrary distribution $\phi(x)$:

$$\mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \mathbb{E}_{x \sim \phi(x)} \frac{\pi(x, \theta)}{\phi(x)} f(x)$$

From importance sampling point of view

Recall Importance Sampling. For arbitrary distribution $\phi(x)$:

$$\mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \mathbb{E}_{x \sim \phi(x)} \frac{\pi(x, \theta)}{\phi(x)} f(x)$$

From importance sampling point of view

Recall Importance Sampling. For arbitrary distribution $\phi(x)$:

$$\mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \mathbb{E}_{x \sim \phi(x)} \frac{\pi(x, \theta)}{\phi(x)} f(x)$$

Let's set $\phi(x) \equiv \pi(x, \theta)$:

From importance sampling point of view

Recall Importance Sampling. For arbitrary distribution $\phi(x)$:

$$\mathbb{E}_{x \sim \pi(x, \theta)} f(x) = \mathbb{E}_{x \sim \phi(x)} \frac{\pi(x, \theta)}{\phi(x)} f(x)$$

Let's set $\phi(x) \equiv \pi(x, \theta)$:

$$\nabla_{\theta} \mathbb{E}_{x \sim \phi(x)} \frac{\pi(x, \theta)}{\phi(x)} f(x) = \mathbb{E}_{x \sim \phi(x)} \frac{\nabla_{\theta} \pi(x, \theta)}{\phi(x)} f(x)$$

Note: that is the same gradient as with log-derivative trick².

²really? Could it even happen otherwise?

Let's apply log-derivative trick to our goal!

$$\nabla_{\theta} \mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R = \mathbb{E}_{\pi(\theta)} \nabla_{\theta} \log \pi(\theta) \mathbb{E}_{\mathcal{T}} R$$

Let's apply log-derivative trick to our goal!

$$\nabla_{\theta} \mathbb{E}_{\pi(\theta)} \mathbb{E}_{\mathcal{T}} R = \mathbb{E}_{\pi(\theta)} \nabla_{\theta} \log \pi(\theta) \mathbb{E}_{\mathcal{T}} R \approx$$

We can estimate this gradient using Monte-Carlo by playing, let's say, one game:

$$\approx \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) R$$

Problems of REINFORCE

× Doesn't work.

Problems of REINFORCE

- × Doesn't work.
 - **Reason:** *high variance* of Monte-Carlo gradient estimation.

Problems of REINFORCE

- × Doesn't work.
 - **Reason:** *high variance* of Monte-Carlo gradient estimation.
 - you can play more than one game for one gradient step, but that doesn't help much.

Proposition

For arbitrary distribution $\pi(\theta)$:

$$\mathbb{E} \nabla_{\theta} \log \pi(\theta) = \int \nabla_{\theta} \pi(\theta) = \nabla_{\theta} \int \pi(\theta) = \nabla_{\theta} 1 = 0$$

Proposition

For arbitrary distribution $\pi(\theta)$:

$$\mathbb{E} \nabla_{\theta} \log \pi(\theta) = \int \nabla_{\theta} \pi(\theta) = \nabla_{\theta} \int \pi(\theta) = \nabla_{\theta} 1 = 0$$



Adding $\mathbb{E} \nabla_{\theta} \log \pi(\theta) b$ for some b to gradient estimate will not lead to bias, but may change variance.

Theorem

$$b = \frac{\mathbb{E}(\nabla_{\theta} \log \pi(\theta))^2 R}{\mathbb{E}(\nabla_{\theta} \log \pi(\theta))^2}$$

is the baseline which leads to the lowest variance.

Lowest variance baseline

Theorem

$$b = \frac{\mathbb{E}(\nabla_{\theta} \log \pi(\theta))^2 R}{\mathbb{E}(\nabla_{\theta} \log \pi(\theta))^2}$$

is the baseline which leads to the lowest variance.

- * similar to average reward, which is also a good baseline.

Strange thing: our gradient estimate depends on R , which includes reward in the first state $r(s_0)$, where we haven't performed any actions.³

³did we make any mistake?

Strange thing: our gradient estimate depends on R , which includes reward in the first state $r(s_0)$, where we haven't performed any actions.³

Let's untangle our goal:

$$\nabla_{\theta} \mathbb{E}_{p(s_1)} (r(s_1) + \mathbb{E}_{a_1 \sim \pi(s_1, \theta)} \mathbb{E}_{p(s_2 | s_1, a)} [r(s_2) + \dots])$$

³did we make any mistake?

Careful REINFORCE

Strange thing: our gradient estimate depends on R , which includes reward in the first state $r(s_0)$, where we haven't performed any actions.³

Let's untangle our goal:

$$\nabla_{\theta} \mathbb{E}_{p(s_1)} \left(r(s_1) + \mathbb{E}_{a_1 \sim \pi(s_1, \theta)} \mathbb{E}_{p(s_2 | s_1, a)} [r(s_2) + \dots] \right) =$$

After carefully applying log-derivative trick:

$$= \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) \right)$$

³did we make any mistake?

Careful REINFORCE

Strange thing: our gradient estimate depends on R , which includes reward in the first state $r(s_0)$, where we haven't performed any actions.³

Let's untangle our goal:

$$\nabla_{\theta} \mathbb{E}_{p(s_1)} \left(r(s_1) + \mathbb{E}_{a_1 \sim \pi(s_1, \theta)} \mathbb{E}_{p(s_2 | s_1, a)} [r(s_2) + \dots] \right) =$$

After carefully applying log-derivative trick:

$$= \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) \right)$$

✓ that's much better!

³did we make any mistake?

Note: $\sum_{t'=t+1}^T r(s_{t'})$ is estimation of $Q^\pi(s_t, a_t)$!

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) \right)$$

Note: $\sum_{t'=t+1}^T r(s_{t'})$ is estimation of $Q^\pi(s_t, a_t)$!

$$\begin{aligned}\nabla &= \mathbb{E} \sum_t^T \nabla_\theta \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) \right) = \\ &= \mathbb{E} \sum_t^T \nabla_\theta \log \pi(a_t | s_t, \theta) Q^\pi(s_t, a_t)\end{aligned}$$

Actor-critic

Note: $\sum_{t'=t+1}^T r(s_{t'})$ is estimation of $Q^\pi(s_t, a_t)$!

$$\begin{aligned}\nabla &= \mathbb{E} \sum_t^T \nabla_\theta \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) \right) = \\ &= \mathbb{E} \sum_t^T \nabla_\theta \log \pi(a_t | s_t, \theta) Q^\pi(s_t, a_t)\end{aligned}$$



Better estimation of $Q^\pi(s, a)$
should lead to lower variance.

Actor-critic

Note: $\sum_{t'=t+1}^T r(s_{t'})$ is estimation of $Q^\pi(s_t, a_t)$!

$$\begin{aligned}\nabla &= \mathbb{E} \sum_t^T \nabla_\theta \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) \right) = \\ &= \mathbb{E} \sum_t^T \nabla_\theta \log \pi(a_t | s_t, \theta) Q^\pi(s_t, a_t)\end{aligned}$$



Better estimation of $Q^\pi(s, a)$
should lead to lower variance.

- * π is an *actor*
- * estimate of $Q^\pi(s, a)$ is a *critic*

Let's insert some baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) Q^{\pi}(s_t, a_t)$$

Let's insert some baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) (Q^{\pi}(s_t, a_t) - b)$$

Let's insert some baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) (Q^{\pi}(s_t, a_t) - b)$$

* Recall average $Q^{\pi}(s_t, a_t)$ is a good baseline.

Let's insert some baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) (Q^{\pi}(s_t, a_t) - b)$$

- * Recall average $Q^{\pi}(s_t, a_t)$ is a good baseline.
- * Recall $\mathbb{E} Q^{\pi}(s_t, a_t) = V^{\pi}(s_t)$

Let's insert some baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) (Q^{\pi}(s_t, a_t) - b)$$

- * Recall average $Q^{\pi}(s_t, a_t)$ is a good baseline.
- * Recall $\mathbb{E} Q^{\pi}(s_t, a_t) = V^{\pi}(s_t)$
- * Recall definition $A^{\pi}(s_t, a_t) = Q^{\pi}(s_t, a_t) - V^{\pi}(s_t)$

Let's insert some baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) A^{\pi}(s_t, a_t)$$

- * Recall average $Q^{\pi}(s_t, a_t)$ is a good baseline.
- * Recall $\mathbb{E} Q^{\pi}(s_t, a_t) = V^{\pi}(s_t)$
- * Recall definition $A^{\pi}(s_t, a_t) = Q^{\pi}(s_t, a_t) - V^{\pi}(s_t)$



Critic can be a second neural net!



Critic can be a second neural net!

Options:

- * approximate $A^\pi(s, a)$



Critic can be a second neural net!

Options:

- * approximate $A^\pi(s, a)$
- * approximate $Q^\pi(s, a)$ ⁴

⁴can we just use Q-learning for this?



Critic can be a second neural net!

Options:

- * approximate $A^\pi(s, a)$
- * approximate $Q^\pi(s, a)$ ⁴
- * approximate $V^\pi(s)$:

⁴can we just use Q-learning for this?



Critic can be a second neural net!

Options:

- * approximate $A^\pi(s, a)$
- * approximate $Q^\pi(s, a)$ ⁴
- * approximate $V^\pi(s)$:

$$Q^\pi(s_t, a_t) - V^\pi(s_t) \approx r(s_{t+1}) + V^\pi(s_{t+1}) - V^\pi(s_t)$$

⁴can we just use Q-learning for this?



Critic can be a second neural net!

Options:

- * approximate $A^\pi(s, a)$
- * approximate $Q^\pi(s, a)$ ⁴
- * approximate $V^\pi(s)$:

$$Q^\pi(s_t, a_t) - V^\pi(s_t) \approx r(s_{t+1}) + V^\pi(s_{t+1}) - V^\pi(s_t)$$

✓ the least complex one! ⁵

⁴can we just use Q-learning for this?

⁵why?

For given state s we can calculate a target $y = V^\pi(s) \approx \sum_{t'=t+1}^T r(s_{t'})$.

At the end of the game, make a step of gradient descent to teach critic.

For given state s we can calculate a target $y = V^\pi(s) \approx \sum_{t'=t+1}^T r(s_{t'})$.

At the end of the game, make a step of gradient descent to teach critic.

Problem: the batch is highly correlated.

For given state s we can calculate a target $y = V^\pi(s) \approx \sum_{t'=t+1}^T r(s_{t'})$.

At the end of the game, make a step of gradient descent to teach critic.

Problem: the batch is highly correlated.

- * well, play more games.

For given state s we can calculate a target $y = V^\pi(s) \approx \sum_{t'=t+1}^T r(s_{t'})$.

At the end of the game, make a step of gradient descent to teach critic.

Problem: the batch is highly correlated.

- * well, play more games.

- :(for one step of gradient descent, yeah...

For given state s we can calculate a target $y = V^\pi(s) \approx \sum_{t'=t+1}^T r(s_{t'})$.

At the end of the game, make a step of gradient descent to teach critic.

Problem: the batch is highly correlated.

- * well, play more games.

- :(for one step of gradient descent, yeah. . .

Alternative: $y = V^\pi(s) \approx r(s') + V^\pi(s')$

Advantage Actor-Critic (A2C) Algorithm:

- get (s, a, r, s')

Advantage Actor-Critic (A2C) Algorithm:

- get (s, a, r, s')
- update critic $\hat{V}(s)$ using target $r + \hat{V}(s')$

Advantage Actor-Critic (A2C) Algorithm:



- get (s, a, r, s')
- update critic $\hat{V}(s)$ using target $r + \hat{V}(s')$
- evaluate $\hat{A}(s, a) = r + \hat{V}(s') - \hat{V}(s)$

Advantage Actor-Critic (A2C) Algorithm:

- get (s, a, r, s')
- update critic $\hat{V}(s)$ using target $r + \hat{V}(s')$
- evaluate $\hat{A}(s, a) = r + \hat{V}(s') - \hat{V}(s)$
- update policy using estimate of gradient $\nabla_{\theta} \log \pi(a \mid s, \theta) \hat{A}(s, a)$

Dealing with two networks

Option 1: just two neural nets





Option 2: shared feature extractor

Dealing with two networks

Option 1: just two neural nets

× obviously redundant



Option 2: shared feature extractor

× may be unstable

✓ lower variance!

- ✓ lower variance!
- × yet policy gradient estimates are not unbiased anymore!⁶

⁶why?

- ✓ lower variance!
- × yet policy gradient estimates are not unbiased anymore!⁶
- × `batch_size = 1`

⁶why?

- ✓ lower variance!
- × yet policy gradient estimates are not unbiased anymore!⁶
- × `batch_size = 1`
 - ✓ do gradient descent step every N game steps.

⁶why?

- ✓ lower variance!
- × yet policy gradient estimates are not unbiased anymore!⁶
- × `batch_size = 1`
 - ✓ do gradient descent step every N game steps.
 - ✓ play several games in parallel.

⁶why?

- ✓ lower variance!
- × yet policy gradient estimates are not unbiased anymore!
- × `batch_size = 1`
 - ✓ do gradient descent step every N game steps.
 - ✓ play several games in parallel.

Check out [this comic](#) about A2C!

Generalized Advantage Estimation (GAE) (2018)

Playing with Q and V ...

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) (Q^{\pi}(s_t, a_t) - V^{\pi}(s_t))$$

In practice we may use separate approximations for $Q^{\pi}(s_t, a_t)$ and baseline $b = V^{\pi}(s_t)$ and play with different ways to do that:

Playing with Q and V ...

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) (Q^{\pi}(s_t, a_t) - V^{\pi}(s_t))$$

In practice we may use separate approximations for $Q^{\pi}(s_t, a_t)$ and baseline $b = V^{\pi}(s_t)$ and play with different ways to do that:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \psi_t$$

ψ_t	bias	variance
$\sum_t^T r(s_t)$	0	very high
$Q^{\pi}(s_t, a_t)$	tolerant	high
$A^{\pi}(s_t, a_t)$	tolerant	low enough
$\sum_t^T r(s_t) - V^{\pi}(s_t)$	0	low

We may use critic **only** for baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) - V^{\pi}(s_t) \right)$$

Eligibility trace

We may use critic only for baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) - V^{\pi}(s_t) \right)$$

✓ unbiased gradient

Eligibility trace

We may use critic only for baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) - V^{\pi}(s_t) \right)$$

- ✓ unbiased gradient
- × higher variance

Eligibility trace

We may use critic only for baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) - V^{\pi}(s_t) \right)$$

✓ unbiased gradient

× higher variance

Or use a compromise (for simplicity $\gamma = 1$):

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^{t+N} r(s_{t'}) + V^{\pi}(s_{t+N}) - V^{\pi}(s_t) \right)$$

Eligibility trace

We may use critic only for baseline:

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^T r(s_{t'}) - V^{\pi}(s_t) \right)$$

✓ unbiased gradient

× higher variance

Or use a compromise (for simplicity $\gamma = 1$):

$$\nabla = \mathbb{E} \sum_t^T \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t+1}^{t+N} r(s_{t'}) + V^{\pi}(s_{t+N}) - V^{\pi}(s_t) \right)$$

× new hyperparameter N

✓ regulates trade-off between variance and bias

So, for different N we have different advantage estimators.

So, for different N we have different advantage estimators.

Generalized Advantage Estimator (2018):



Create an ensemble out of them!

So, for different N we have different advantage estimators.

Generalized Advantage Estimator (2018):



Create an ensemble out of them!

Let $A_{(N)}^{\pi}(s_t, a_t)$ be a N -step advantage estimator:

$$A_{(N)}^{\pi} = \sum_{t'=t+1}^{t+N} r(s_{t'}) + V^{\pi}(s_{t+N}) - V^{\pi}(s_t)$$

So, for different N we have different advantage estimators.

Generalized Advantage Estimator (2018):



Create an ensemble out of them!

Let $A_{(N)}^{\pi}(s_t, a_t)$ be a N -step advantage estimator:

$$A_{(N)}^{\pi} = \sum_{t'=t+1}^{t+N} r(s_{t'}) + V^{\pi}(s_{t+N}) - V^{\pi}(s_t)$$

Let's take exponentially-weighted average:

$$A_{(\text{GAE})}^{\pi}(s_t, a_t) = (1 - \lambda)(A_{(1)}^{\pi} + \lambda A_{(2)}^{\pi} + \lambda^2 A_{(3)}^{\pi} + \dots)$$

Move convenient formula:

$$A_{(\text{GAE})}^{\pi}(s_t, a_t) = \sum_{i=0}^{\infty} (\lambda \gamma)^i (r(s_{t+i}) + \gamma V^{\pi}(s_{t+i+1}) - V^{\pi}(s_{t+i}))$$

Move convenient formula:

$$A_{(\text{GAE})}^{\pi}(s_t, a_t) = \sum_{i=0}^{\infty} (\lambda \gamma)^i (r(s_{t+i}) + \gamma V^{\pi}(s_{t+i+1}) - V^{\pi}(s_{t+i}))$$

- * $\lambda = 0$: A2C algorithm
- * $\lambda = 1$: infinite eligibility trace algorithm

Move convenient formula:

$$A_{(\text{GAE})}^{\pi}(s_t, a_t) = \sum_{i=0}^{\infty} (\lambda \gamma)^i (r(s_{t+i}) + \gamma V^{\pi}(s_{t+i+1}) - V^{\pi}(s_{t+i}))$$

- * $\lambda = 0$: A2C algorithm
- * $\lambda = 1$: infinite eligibility trace algorithm
- * the balance is in between...



NEXT: TRPO

Trust Region Policy Optimization (TRPO) (2017)

Problem: Actor-Critic algorithm is *on-policy*.

- × we **throw away** obtained data after one optimization step.

Problem: Actor-Critic algorithm is *on-policy*.

× we throw away obtained data after one optimization step.

! but we have to do this!

$$\nabla(\theta) = \mathbb{E}_{\mathcal{T} \sim \pi(\theta)} \sum_t^T \nabla_{\theta} \log \pi(a_t \mid s_t, \theta) A^{\pi}(s_t, a_t)$$

Problem: Actor-Critic algorithm is *on-policy*.

× we **throw away** obtained data after one optimization step.

! but we have to do this!

$$\nabla(\theta) = \mathbb{E}_{\mathcal{T} \sim \pi(\theta)} \sum_t^T \nabla_{\theta} \log \pi(a_t \mid s_t, \theta) A^{\pi}(s_t, a_t)$$



Use importance sampling!

Let denote $P(\mathcal{T} \mid \pi)$ a probability of trajectory under policy π :

$$P(\mathcal{T} \mid \pi) = p(s_0) \prod_{t=0} [\pi(a_t \mid s_t) p(s_{t+1} \mid s_t, a_t)]$$

Off-policy Actor-Critic

Let denote $P(\mathcal{T} \mid \pi)$ a probability of trajectory under policy π :

$$P(\mathcal{T} \mid \pi) = p(s_0) \prod_{t=0} [\pi(a_t \mid s_t) p(s_{t+1} \mid s_t, a_t)]$$

Then off-policy actor-critic gradient estimation can be obtained:

$$\nabla(\theta) = \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\frac{P(\mathcal{T} \mid \pi)}{P(\mathcal{T} \mid \tilde{\pi})} \sum_t^T \nabla_{\theta} \log \pi(a_t \mid s_t, \theta) A^{\pi}(s_t, a_t) \right]$$

Off-policy Actor-Critic

Let denote $P(\mathcal{T} \mid \pi)$ a probability of trajectory under policy π :

$$P(\mathcal{T} \mid \pi) = p(s_0) \prod_{t=0} [\pi(a_t \mid s_t) p(s_{t+1} \mid s_t, a_t)]$$

Then off-policy actor-critic gradient estimation can be obtained:

$$\nabla(\theta) = \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\frac{P(\mathcal{T} \mid \pi)}{P(\mathcal{T} \mid \tilde{\pi})} \sum_t \nabla_{\theta} \log \pi(a_t \mid s_t, \theta) A^{\pi}(s_t, a_t) \right]$$

- × though transition probability reduce, this *importance sampling weight* tends to be very close to 0.



May be if π is close to $\tilde{\pi}$, this weight is practically acceptable



May be if π is close to $\tilde{\pi}$, this weight is practically acceptable

Trust-Region Policy Optimization (2017) hints:

- a lot of theory on relative performance of two close policies
- attempt to build policy optimization procedure with guarantees of optimizing the objective.⁶
- practical application of **natural policy gradients**.

⁶what is an obvious drawback of procedure with such property?



May be if π is close to $\tilde{\pi}$, this weight is practically acceptable

Trust-Region Policy Optimization (2017) hints:

- a lot of theory on relative performance of two close policies
- attempt to build policy optimization procedure with guarantees of optimizing the objective.⁶
- practical application of **natural policy gradients**.
- × doesn't provide enthusiastic results on practice...

⁶what is an obvious drawback of procedure with such property?

Relative Policy Performance Identity

Let's denote $J(\pi)$ a performance of policy π , i.e. our objective:

$$J(\pi) \stackrel{\text{def}}{=} \mathbb{E}_{\mathcal{T} \sim \pi} \sum_{t=0} \gamma^t r(s_t) = \mathbb{E}_{s_0} V^\pi(s_0)$$

Relative Policy Performance Identity

Let's denote $J(\pi)$ a performance of policy π , i.e. our objective:

$$J(\pi) \stackrel{\text{def}}{=} \mathbb{E}_{\mathcal{T} \sim \pi} \sum_{t=0} \gamma^t r(s_t) = \mathbb{E}_{s_0} V^\pi(s_0)$$

Theorem (Kakade & Langford, 2002):

$$J(\tilde{\pi}) - J(\pi) = \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t A^\pi(s_t, a_t)$$

Relative Policy Performance Identity: Proof

$$J(\tilde{\pi}) - J(\pi) = \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - J(\pi) =$$

Relative Policy Performance Identity: Proof

$$\begin{aligned} J(\tilde{\pi}) - J(\pi) &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - J(\pi) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - \mathbb{E}_{s_0} V^{\pi}(s_0) = \end{aligned}$$

Relative Policy Performance Identity: Proof

$$\begin{aligned} J(\tilde{\pi}) - J(\pi) &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0}^{\infty} \gamma^t r(s_t) - J(\pi) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0}^{\infty} \gamma^t r(s_t) - \mathbb{E}_{s_0} V^{\pi}(s_0) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t) - V^{\pi}(s_0) \right] = \end{aligned}$$

Relative Policy Performance Identity: Proof

$$\begin{aligned} J(\tilde{\pi}) - J(\pi) &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - J(\pi) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - \mathbb{E}_{s_0} V^{\pi}(s_0) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\sum_{t=0} \gamma^t r(s_t) - V^{\pi}(s_0) \right] = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\sum_{t=0} \gamma^t r(s_t) + \sum_{t=0} [\gamma^{t+1} V^{\pi}(s_{t+1}) - \gamma^t V^{\pi}(s_t)] \right] = \end{aligned}$$

Relative Policy Performance Identity: Proof

$$\begin{aligned} J(\tilde{\pi}) - J(\pi) &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - J(\pi) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - \mathbb{E}_{s_0} V^{\pi}(s_0) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\sum_{t=0} \gamma^t r(s_t) - V^{\pi}(s_0) \right] = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\sum_{t=0} \gamma^t r(s_t) + \sum_{t=0} [\gamma^{t+1} V^{\pi}(s_{t+1}) - \gamma^t V^{\pi}(s_t)] \right] = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t (r(s_t) + \gamma V^{\pi}(s_{t+1}) - V^{\pi}(s_t)) = \end{aligned}$$

Relative Policy Performance Identity: Proof

$$\begin{aligned} J(\tilde{\pi}) - J(\pi) &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - J(\pi) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t r(s_t) - \mathbb{E}_{s_0} V^{\pi}(s_0) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\sum_{t=0} \gamma^t r(s_t) - V^{\pi}(s_0) \right] = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \left[\sum_{t=0} \gamma^t r(s_t) + \sum_{t=0} [\gamma^{t+1} V^{\pi}(s_{t+1}) - \gamma^t V^{\pi}(s_t)] \right] = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t (r(s_t) + \gamma V^{\pi}(s_{t+1}) - V^{\pi}(s_t)) = \\ &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t A^{\pi}(s_t, a_t) \end{aligned}$$

Applying importance sampling

Denote $d_\pi(s)$ a *discounted state-visitation probability* for policy π :

$$d_\pi(s) = (1 - \gamma) \sum_{t=0} \gamma^t \mathcal{P}(s_t = s)$$

Applying importance sampling

Denote $d_\pi(s)$ a *discounted state-visitation probability* for policy π :

$$d_\pi(s) = (1 - \gamma) \sum_{t=0} \gamma^t \mathcal{P}(s_t = s)$$

Let's separate the expectation over trajectory to the expectation over policy choices and over state transitions:

$$J(\tilde{\pi}) - J(\pi) = \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t A^\pi(s_t, a_t) =$$

Applying importance sampling

Denote $d_\pi(s)$ a *discounted state-visitation probability* for policy π :

$$d_\pi(s) = (1 - \gamma) \sum_{t=0} \gamma^t \mathcal{P}(s_t = s)$$

Let's separate the expectation over trajectory to the expectation over policy choices and over state transitions:

$$\begin{aligned} J(\tilde{\pi}) - J(\pi) &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t A^\pi(s_t, a_t) = \\ &= \frac{1}{1 - \gamma} \mathbb{E}_{s \sim d_\pi} \mathbb{E}_{a \sim \tilde{\pi}} A^\pi(s_t, a_t) = \end{aligned}$$

Applying importance sampling

Denote $d_\pi(s)$ a *discounted state-visitation probability* for policy π :

$$d_\pi(s) = (1 - \gamma) \sum_{t=0} \gamma^t \mathcal{P}(s_t = s)$$

Let's separate the expectation over trajectory to the expectation over policy choices and over state transitions:

$$\begin{aligned} J(\tilde{\pi}) - J(\pi) &= \mathbb{E}_{\mathcal{T} \sim \tilde{\pi}} \sum_{t=0} \gamma^t A^\pi(s_t, a_t) = \\ &= \frac{1}{1 - \gamma} \mathbb{E}_{s \sim d_{\tilde{\pi}}} \mathbb{E}_{a \sim \tilde{\pi}} A^\pi(s_t, a_t) = \\ \{\text{importance sampling}\} &= \frac{1}{1 - \gamma} \mathbb{E}_{s \sim d_{\tilde{\pi}}} \mathbb{E}_{a \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^\pi(s_t, a_t) \end{aligned}$$

Approximation

Suppose π is current policy and $\tilde{\pi}$ is a policy after one optimization step. To make this step, we can't sample from $d_{\tilde{\pi}}$.

Approximation

Suppose π is current policy and $\tilde{\pi}$ is a policy after one optimization step. To make this step, we can't sample from $d_{\tilde{\pi}}$.

Available approximation:

$$\begin{aligned} & \frac{1}{1-\gamma} \mathbb{E}_{s \sim d_{\tilde{\pi}}} \mathbb{E}_{a \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \approx \\ & \approx \frac{1}{1-\gamma} \mathbb{E}_{s \sim d_{\pi}} \mathbb{E}_{a \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) = \end{aligned}$$

Approximation

Suppose π is current policy and $\tilde{\pi}$ is a policy after one optimization step. To make this step, we can't sample from $d_{\tilde{\pi}}$.

Available approximation:

$$\begin{aligned} & \frac{1}{1-\gamma} \mathbb{E}_{s \sim d_{\tilde{\pi}}} \mathbb{E}_{a \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \approx \\ & \approx \frac{1}{1-\gamma} \mathbb{E}_{s \sim d_{\pi}} \mathbb{E}_{a \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) = \\ & = \mathbb{E}_{\mathcal{T} \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \stackrel{\text{def}}{=} L(\tilde{\pi}) \end{aligned}$$

Approximation

Suppose π is current policy and $\tilde{\pi}$ is a policy after one optimization step. To make this step, we can't sample from $d_{\tilde{\pi}}$.

Available approximation:

$$\begin{aligned} & \frac{1}{1-\gamma} \mathbb{E}_{s \sim d_{\tilde{\pi}}} \mathbb{E}_{a \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \approx \\ & \approx \frac{1}{1-\gamma} \mathbb{E}_{s \sim d_{\pi}} \mathbb{E}_{a \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) = \\ & = \mathbb{E}_{\mathcal{T} \sim \pi} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \stackrel{\text{def}}{=} L(\tilde{\pi}) \end{aligned}$$

Theorem:

If ε is the approximation error:

$$|\varepsilon| \leq \text{Const } KL^{\max}(\tilde{\pi} \parallel \pi)$$

The familiar gradients...

Let π 's parameters be θ_k (fixed), $\tilde{\pi}$'s parameters be θ .

To optimize θ , let's find $\nabla_{\theta} L(\tilde{\pi}(\theta))|_{\theta_k}$:

$$\nabla_{\theta} L(\tilde{\pi}(\theta))|_{\theta_k} = \nabla_{\theta} \left[\mathbb{E}_{\mathcal{T} \sim \pi} \frac{\tilde{\pi}_{\theta}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \right] \Big|_{\theta_k} =$$

The familiar gradients...

Let π 's parameters be θ_k (fixed), $\tilde{\pi}$'s parameters be θ .

To optimize θ , let's find $\nabla_{\theta} L(\tilde{\pi}(\theta))|_{\theta_k}$:

$$\begin{aligned}\nabla_{\theta} L(\tilde{\pi}(\theta))|_{\theta_k} &= \nabla_{\theta} \left[\mathbb{E}_{\mathcal{T} \sim \pi} \frac{\tilde{\pi}_{\theta}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \right] \Big|_{\theta_k} = \\ &= \mathbb{E}_{\mathcal{T} \sim \pi} \frac{\nabla_{\theta} \tilde{\pi}_{\theta}(a_t | s_t)|_{\theta_k}}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) =\end{aligned}$$

The familiar gradients...

Let π 's parameters be θ_k (fixed), $\tilde{\pi}$'s parameters be θ .

To optimize θ , let's find $\nabla_{\theta} L(\tilde{\pi}(\theta))|_{\theta_k}$:

$$\begin{aligned}\nabla_{\theta} L(\tilde{\pi}(\theta))|_{\theta_k} &= \nabla_{\theta} \left[\mathbb{E}_{\mathcal{T} \sim \pi} \frac{\tilde{\pi}_{\theta}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \right] \Big|_{\theta_k} = \\ &= \mathbb{E}_{\mathcal{T} \sim \pi} \frac{\nabla_{\theta} \tilde{\pi}_{\theta}(a_t | s_t)|_{\theta_k}}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) = \\ \{\pi \equiv \tilde{\pi}(\theta_k)\} &= \mathbb{E}_{\mathcal{T} \sim \pi} \nabla_{\theta} \log \tilde{\pi}_{\theta}(a_t | s_t)|_{\theta_k} A^{\pi}(s_t, a_t)\end{aligned}$$

The familiar gradients...

Let π 's parameters be θ_k (fixed), $\tilde{\pi}$'s parameters be θ .

To optimize θ , let's find $\nabla_{\theta} L(\tilde{\pi}(\theta))|_{\theta_k}$:

$$\begin{aligned}\nabla_{\theta} L(\tilde{\pi}(\theta))|_{\theta_k} &= \nabla_{\theta} \left[\mathbb{E}_{\mathcal{T} \sim \pi} \frac{\tilde{\pi}_{\theta}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \right] \Big|_{\theta_k} = \\ &= \mathbb{E}_{\mathcal{T} \sim \pi} \frac{\nabla_{\theta} \tilde{\pi}_{\theta}(a_t | s_t)|_{\theta_k}}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) = \\ \{\pi \equiv \tilde{\pi}(\theta_k)\} &= \mathbb{E}_{\mathcal{T} \sim \pi} \nabla_{\theta} \log \tilde{\pi}_{\theta}(a_t | s_t)|_{\theta_k} A^{\pi}(s_t, a_t)\end{aligned}$$



From approximation error bound follows:

$$J(\tilde{\pi}) - J(\pi) \geq L(\tilde{\pi}) - C KL^{\max}(\tilde{\pi} \parallel \pi)$$

Improvement guarantees

From approximation error bound follows:

$$J(\tilde{\pi}) - J(\pi) \geq L(\tilde{\pi}) - C KL^{\max}(\tilde{\pi} \parallel \pi)$$

Consider the lower bound optimization procedure:

$$\pi_{k+1} = \underset{\tilde{\pi}}{\operatorname{argmax}} [L(\tilde{\pi}) - C KL^{\max}(\tilde{\pi} \parallel \pi_k)]$$

Improvement guarantees

From approximation error bound follows:

$$J(\tilde{\pi}) - J(\pi) \geq L(\tilde{\pi}) - C KL^{\max}(\tilde{\pi} \parallel \pi)$$

Consider the lower bound optimization procedure:

$$\pi_{k+1} = \underset{\tilde{\pi}}{\operatorname{argmax}} [L(\tilde{\pi}) - C KL^{\max}(\tilde{\pi} \parallel \pi_k)]$$

Then:

$$\begin{aligned} J(\pi_{k+1}) - J(\pi_k) &\geq L(\pi_{k+1}) - C KL^{\max}(\pi_{k+1} \parallel \pi_k) \geq \\ &\geq L(\pi_k) - C KL^{\max}(\pi_k \parallel \pi_k) = 0 - 0 = 0 \end{aligned}$$

Improvement guarantees

From approximation error bound follows:

$$J(\tilde{\pi}) - J(\pi) \geq L(\tilde{\pi}) - C KL^{\max}(\tilde{\pi} \parallel \pi)$$

Consider the lower bound optimization procedure:

$$\pi_{k+1} = \underset{\tilde{\pi}}{\operatorname{argmax}} [L(\tilde{\pi}) - C KL^{\max}(\tilde{\pi} \parallel \pi_k)]$$

Then:

$$\begin{aligned} J(\pi_{k+1}) - J(\pi_k) &\geq L(\pi_{k+1}) - C KL^{\max}(\pi_{k+1} \parallel \pi_k) \geq \\ &\geq L(\pi_k) - C KL^{\max}(\pi_k \parallel \pi_k) = 0 - 0 = 0 \end{aligned}$$

✓ procedure guarantees to improve $J(\pi)$!

$$\pi_{k+1} = \underset{\tilde{\pi}}{\operatorname{argmax}} \left[\mathbb{E}_{\mathcal{T} \sim \pi_k} \frac{\tilde{\pi}(a_t | s_t)}{\pi_k(a_t | s_t)} A^{\pi_k}(s_t, a_t) - C \operatorname{KL}^{\max}(\tilde{\pi} \parallel \pi_k) \right]$$

- × A^{π} is never precise.
- × expectations estimations are never precise.

$$\pi_{k+1} = \underset{\tilde{\pi}}{\operatorname{argmax}} \left[\mathbb{E}_{\mathcal{T} \sim \pi_k} \frac{\tilde{\pi}(a_t | s_t)}{\pi_k(a_t | s_t)} A^{\pi_k}(s_t, a_t) - C \text{KL}^{\max}(\tilde{\pi} \parallel \pi_k) \right]$$

- × A^{π} is never precise.
- × expectations estimations are never precise.
- × we can't calculate maximal divergence between two policies over *all* states.

$$\pi_{k+1} = \underset{\tilde{\pi}}{\operatorname{argmax}} \left[\mathbb{E}_{\mathcal{T} \sim \pi_k} \frac{\tilde{\pi}(a_t | s_t)}{\pi_k(a_t | s_t)} A^{\pi_k}(s_t, a_t) - C KL^{\max}(\tilde{\pi} \parallel \pi_k) \right]$$

- × A^{π} is never precise.
- × expectations estimations are never precise.
- × we can't calculate maximal divergence between two policies over *all* states.



TRPO: $KL^{\max}(\tilde{\pi} \parallel \pi_k) \approx \mathbb{E}_{s \sim d_{\pi_k}} KL(\tilde{\pi} \parallel \pi_k)[s]$

$$\pi_{k+1} = \underset{\tilde{\pi}}{\operatorname{argmax}} \left[\mathbb{E}_{\mathcal{T} \sim \pi_k} \frac{\tilde{\pi}(a_t | s_t)}{\pi_k(a_t | s_t)} A^{\pi_k}(s_t, a_t) - \textcolor{brown}{C} KL^{\max}(\tilde{\pi} \parallel \pi_k) \right]$$

- × A^{π} is never precise.
- × expectations estimations are never precise.
- × we can't calculate maximal divergence between two policies over *all* states.



TRPO: $KL^{\max}(\tilde{\pi} \parallel \pi_k) \approx \mathbb{E}_{s \sim d_{\pi_k}} KL(\tilde{\pi} \parallel \pi_k)[s]$

- × the constant over here is huge when γ is close to 1 and depends on MDP characteristics.

$$\pi_{k+1} = \underset{\tilde{\pi}}{\operatorname{argmax}} \left[\mathbb{E}_{\mathcal{T} \sim \pi_k} \frac{\tilde{\pi}(a_t | s_t)}{\pi_k(a_t | s_t)} A^{\pi_k}(s_t, a_t) - C KL^{\max}(\tilde{\pi} \parallel \pi_k) \right]$$

- × A^{π} is never precise.
- × expectations estimations are never precise.
- × we can't calculate maximal divergence between two policies over *all* states.



TRPO: $KL^{\max}(\tilde{\pi} \parallel \pi_k) \approx \mathbb{E}_{s \sim d_{\pi_k}} KL(\tilde{\pi} \parallel \pi_k)[s]$

- × the constant over here is huge when γ is close to 1 and depends on MDP characteristics.




TRPO: Trust-Region optimization scheme!

$$\begin{cases} \pi_{k+1} = \mathbb{E}_{\mathcal{T} \sim \pi_k} \frac{\tilde{\pi}(a_t|s_t)}{\pi_k(a_t|s_t)} A^{\pi_k}(s_t, a_t) \rightarrow \max_{\tilde{\pi}} \\ \text{s.t. } \mathbb{E}_{s \sim d_{\pi_k}} KL(\tilde{\pi} \parallel \pi_k)[s] \leq \delta \end{cases}$$


Trust-Region optimization

$$\begin{cases} \pi_{k+1} = \mathbb{E}_{\mathcal{T} \sim \pi_k} \frac{\tilde{\pi}(a_t|s_t)}{\pi_k(a_t|s_t)} A^{\pi_k}(s_t, a_t) \rightarrow \max_{\tilde{\pi}} \\ \text{s.t. } \mathbb{E}_{s \sim d_{\pi_k}} KL(\tilde{\pi} \parallel \pi_k)[s] \leq \delta \end{cases}$$

× δ is a hyperparameter. 

Trust-Region optimization

$$\begin{cases} \pi_{k+1} = \mathbb{E}_{\mathcal{T} \sim \pi_k} \frac{\tilde{\pi}(a_t|s_t)}{\pi_k(a_t|s_t)} A^{\pi_k}(s_t, a_t) \rightarrow \max_{\tilde{\pi}} \\ \text{s.t. } \mathbb{E}_{s \sim d_{\pi_k}} KL(\tilde{\pi} \parallel \pi_k)[s] \leq \delta \end{cases}$$

- × δ is a hyperparameter. 
- ✓ respects distance in policy space!
 - also known in theory as *natural gradient*. In previous policy gradient methods we implicitly used the constrain

$$\|\tilde{\theta} - \theta_k\|_2^2 \leq \alpha$$

where α was learning rate of optimizer.

Natural Policy Gradient

Metric in most general form may depend from current coordinates:

$$\rho(x, x + d) = d^T G(x) d$$

- $G(x)$ is called *metric tensor*.

Natural Policy Gradient

Metric in most general form may depend from current coordinates:

$$\rho(x, x + d) = d^T G(x) d$$

- $G(x)$ is called *metric tensor*.

Theorem:

For space of policies, *Fisher information matrix* is metric tensor:

$$H(\theta) = \mathbb{E}_{a \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a | s) \log \pi_\theta(a | s)^T]$$

Natural Policy Gradient

Metric in most general form may depend from current coordinates:

$$\rho(x, x + d) = d^T G(x) d$$

- $G(x)$ is called *metric tensor*.

Theorem:

For space of policies, *Fisher information matrix* is metric tensor:

$$H(\theta) = \mathbb{E}_{a \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a | s) \log \pi_\theta(a | s)^T]$$

Main natural gradient property (parametrization invariance)

For any parametrization π_θ

$$H^{-1} \nabla_\theta \pi_\theta$$

is the same vector in policies space.

Recalling standard optimization methods to solve constraint task:

$$\begin{cases} L(\theta) \rightarrow \max_{\theta} \\ \text{s.t.} \quad \mathbb{E}_{s \sim d_{\pi(\theta_k)}} KL(\pi(\theta) \parallel \pi(\theta_k))[s] \leq \delta \end{cases}$$

Practical application

Recalling standard optimization methods to solve constraint task:

$$\begin{cases} L(\theta) \rightarrow \max_{\theta} \\ \text{s.t. } \mathbb{E}_{s \sim d_{\pi(\theta_k)}} KL(\pi(\theta) \parallel \pi(\theta_k))[s] \leq \delta \end{cases}$$

Linear approximation of optimized objective:

$$L(\theta) \approx L(\theta_k) + g^T(\theta - \theta_k) \quad \text{where } g = \nabla_{\theta} L(\theta)|_{\theta_k}$$

Practical application

Recalling standard optimization methods to solve constraint task:

$$\begin{cases} L(\theta) \rightarrow \max_{\theta} \\ \text{s.t.} \quad \mathbb{E}_{s \sim d_{\pi(\theta_k)}} KL(\pi(\theta) \parallel \pi(\theta_k))[s] \leq \delta \end{cases}$$

Linear approximation of optimized objective:

$$L(\theta) \approx L(\theta_k) + g^T(\theta - \theta_k) \quad \text{where} \quad g = \nabla_{\theta} L(\theta)|_{\theta_k}$$

Quadratic approximation of constraint ⁷:

$$\begin{aligned} \mathbb{E}_s KL(\pi(\theta) \parallel \pi(\theta_k))[s] &\approx (\theta - \theta_k)^T H (\theta - \theta_k) \\ \text{where} \quad H &= \mathbb{E}_s \nabla_{\theta}^2 KL(\pi(\theta) \parallel \pi(\theta_k))[s]|_{\theta_k} \end{aligned}$$

⁷where is linear term?

Trust-Region optimization procedure

Theorem:

$\nabla_{\theta}^2 KL(\pi(\theta) \parallel \pi(\theta_k))[s]|_{\theta_k}$ is Fisher information matrix.

- ✓ that's why solving this task is equivalent to gradient ascent with natural policy gradient!

Trust-Region optimization procedure

Theorem:

$\nabla_{\theta}^2 KL(\pi(\theta) \parallel \pi(\theta_k))[s]|_{\theta_k}$ is Fisher information matrix.

- ✓ that's why solving this task is equivalent to gradient ascent with natural policy gradient!

Solution (derived with K.K.T. theorem):

$$\theta_{k+1} = \theta_k + \sqrt{\frac{2\delta}{g_k^T H_k^{-1} g_k}} H_k^{-1} g_k$$

Trust-Region optimization procedure

Theorem:

$\nabla_{\theta}^2 KL(\pi(\theta) \parallel \pi(\theta_k))[s]|_{\theta_k}$ is Fisher information matrix.

- ✓ that's why solving this task is equivalent to gradient ascent with natural policy gradient!

Solution (derived with K.K.T. theorem):

$$\theta_{k+1} = \theta_k + \sqrt{\frac{2\delta}{g_k^T H_k^{-1} g_k}} H_k^{-1} g_k$$

- ✓ δ substitutes learning rate.
- × g_k, H_k can only be estimated via samples.



Trust-Region optimization procedure


Theorem:

$\nabla_{\theta}^2 KL(\pi(\theta) \parallel \pi(\theta_k))[s]|_{\theta_k}$ is Fisher information matrix.

- ✓ that's why solving this task is equivalent to gradient ascent with natural policy gradient!

Solution (derived with K.K.T. theorem):

$$\theta_{k+1} = \theta_k + \sqrt{\frac{2\delta}{g_k^T H_k^{-1} g_k}} H_k^{-1} g_k$$

- ✓ δ substitutes learning rate.
- × g_k, H_k can only be estimated via samples. 
- × **Problem:** how to compute H_k^{-1} on practice? For neural networks with N parameters inversion complexity is $\mathcal{O}(N^3)!$..

Conjugate Gradients saving the day

Remembering CG algorithm:

- solves system of linear equations $H_k x = g_k$.

Conjugate Gradients saving the day

Remembering CG algorithm:

- solves system of linear equations $H_k x = g_k$.
- ✓ after j iterations returns sub-optimal solution (approximation of $H^{-1}g$, optimal in Krylov subspace, $\mathcal{L}(g, Hg, H^2g \dots H^{j-1}g)$)

Conjugate Gradients saving the day

Remembering CG algorithm:

- solves system of linear equations $H_k x = g_k$.
- ✓ after j iterations returns sub-optimal solution (approximation of $H^{-1}g$, optimal in Krylov subspace, $\mathcal{L}(g, Hg, H^2g \dots H^{j-1}g)$)
- ✓ only $f(v) = H_k v$ is required.
 - = can be implemented on PyTorch!

Conjugate Gradients saving the day

Remembering CG algorithm:

- solves system of linear equations $H_k x = g_k$.
- ✓ after j iterations returns sub-optimal solution (approximation of $H^{-1}g$, optimal in Krylov subspace, $\mathcal{L}(g, Hg, H^2g \dots H^{j-1}g)$)
- ✓ only $f(v) = H_k v$ is required.
 - = can be implemented on PyTorch!






With all these approximations no theoretical guarantees remain, of course.



With all these approximations no theoretical guarantees remain, of course.

What is suggested in different papers?

- NPG (2005): just use $H^{-1}g$ as gradient, computed somehow, without looking for improvement guarantees.



With all these approximations no theoretical guarantees remain, of course.

What is suggested in different papers?

- NPG (2005): just use $H^{-1}g$ as gradient, computed somehow, without looking for improvement guarantees.
- TRPO (2017): use line search with basic backtracking to guarantee $L(\pi) \geq 0$ and $KL(\pi \parallel \pi_k) < \delta$.



With all these approximations no theoretical guarantees remain, of course.

What is suggested in different papers?

- NPG (2005): just use $H^{-1}g$ as gradient, computed somehow, without looking for improvement guarantees.
- TRPO (2017): use line search with basic backtracking to guarantee $L(\pi) \geq 0$ and $KL(\pi \parallel \pi_k) < \delta$.
- PPO (2017): coming soon.
- ACKTR (2017): coming soon.